



Rendering Worlds with Two Triangles with raytracing on the GPU in 4096 bytes

Iñigo Quilez – iq/rgba

August 22 at NVSCENE 08



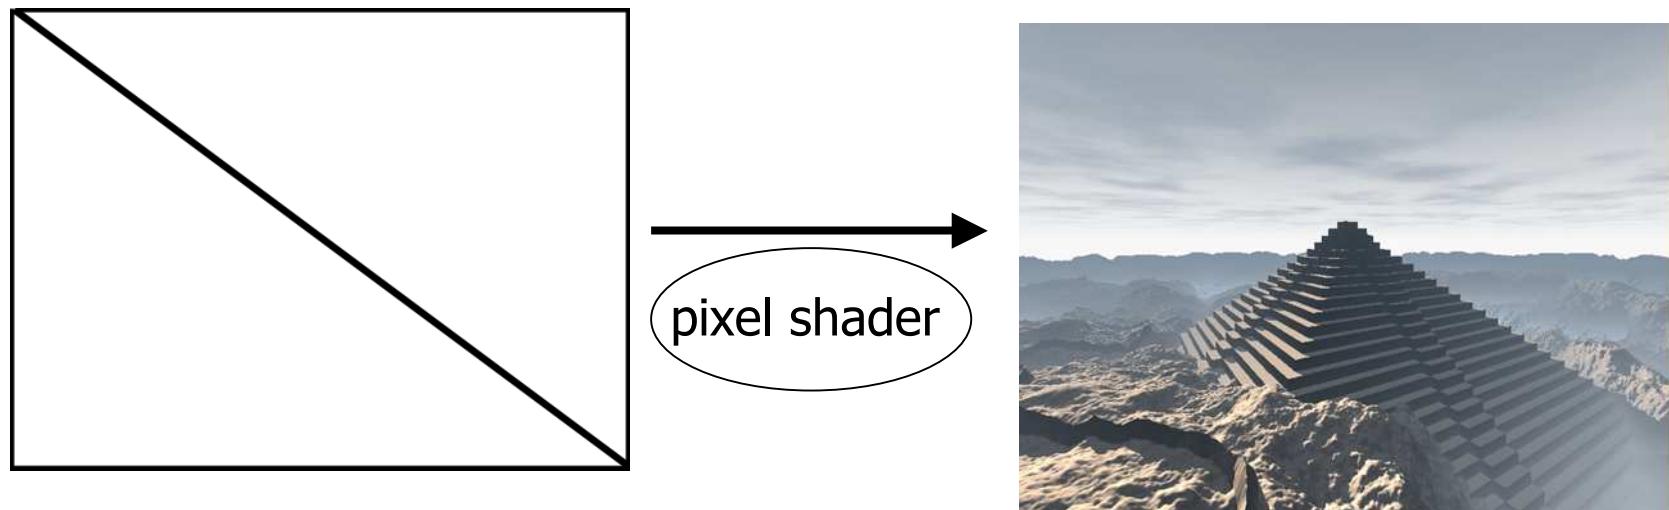
A bit of context

- Amazing progression in raw GPU power.
- Shaders 3 and 4 flexible enough for
 - Experimenting with new techniques.
 - Revival of some old-school effects (at a higher quality than ever).
- Unexpected benefits:
 - Easy to set up and very compact code.
 - 4k demo coders have jumped into it.



A bit of context

- The idea: draw two triangles that cover the entire screen area, and invoke a pixel shader that will create an animated or static image.



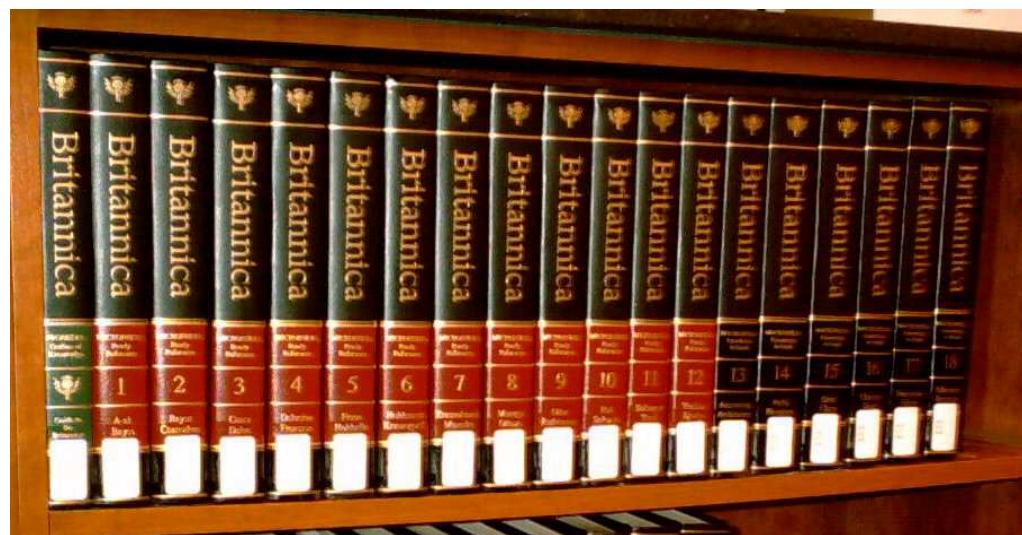
- Make the complete demo self-contained in no more 4096 bytes (that includes the “engine”, music, shaders, animations, textures and everything).

A bit of context

- How much is a kilobyte ?



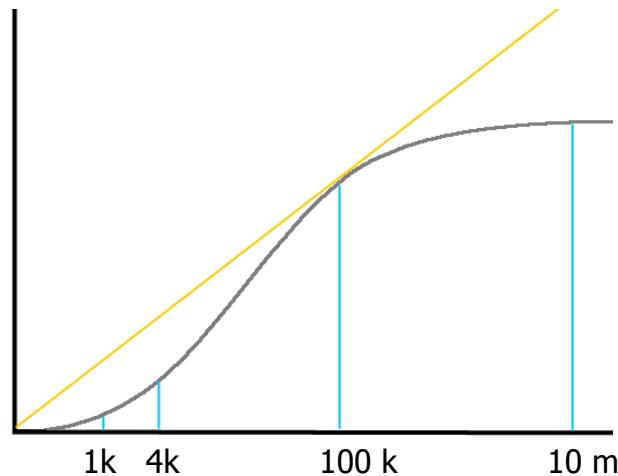
This is the size of a 4 kbytes production



This is the size of a 64 megabytes demo or video

A bit of context

- Probably not a fair comparison (we cannot blame demo coders for being lazy compared to intro coders).
- The “visual_beauty” is not a linear function of the size in kilobytes.



- Speculation: With current technology, the optimal “vibes per kilobyte” (aka result/effort ratio, or “wow factor”) is around 100 kb productions.



Index

- Old-school effects are back
- Rendering with distance fields



Index

- Old-school effects are back
- Rendering with distance fields



Old-school effects are back

- Filling the screen with a shader, and producing an image or animation from it, only works for algorithms and effects that follow this pattern:

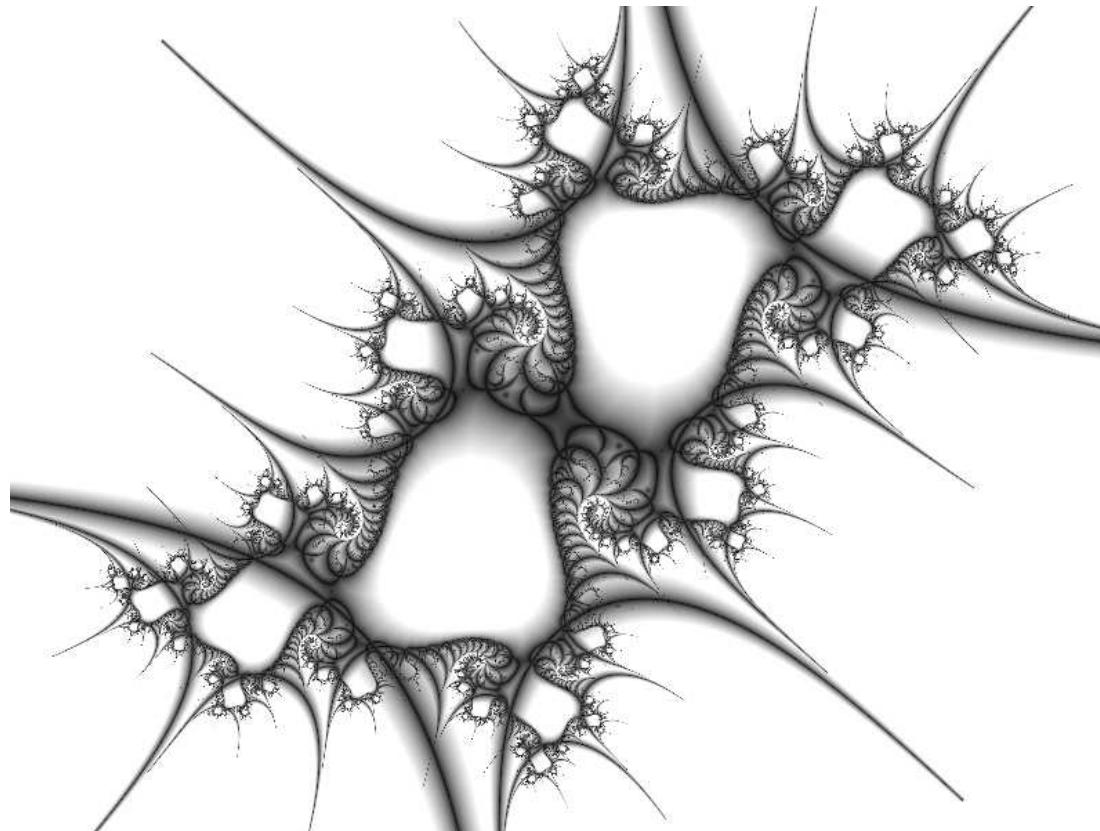
```
for( each pixel p )  
{  
    outputColor = doSomething( p );  
}
```

- This doesn't naturally extend to effects that need to do operations across pixels (gather and scatter operations). Multipass techniques can be used, but
 - it might actually be slower than doing it on the CPU
 - it's not elegant
 - it's not very compact for 4k demos



Old-school effects are back

- Julia and Mandelbrot sets (the “hello world” of gfx programming)





Old-school effects are back

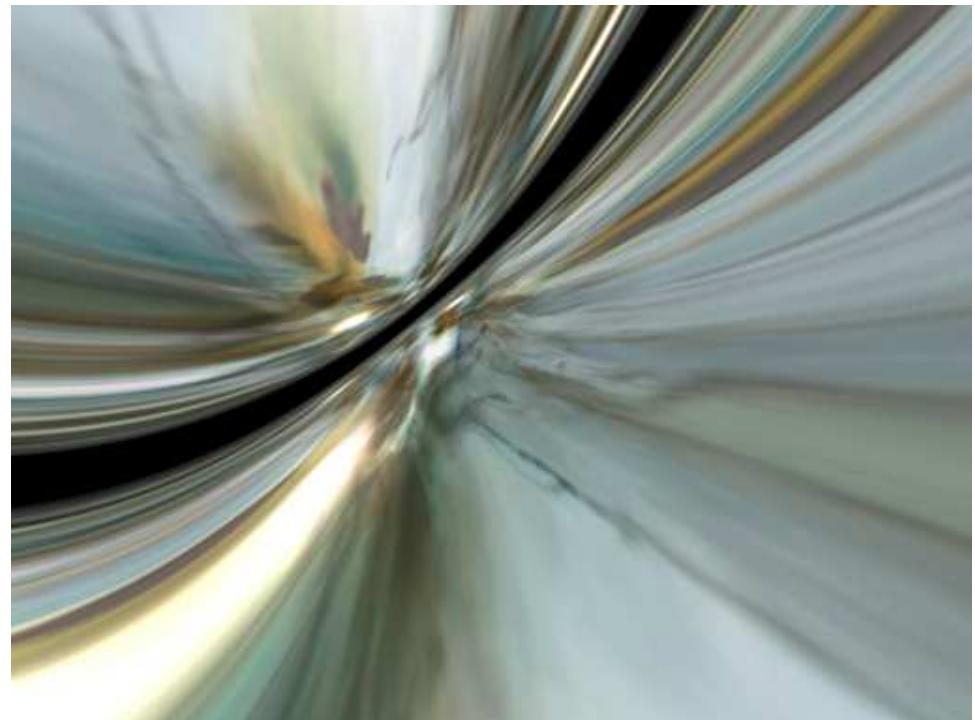
- Plane deformations

Oldschoold software version

```
for( int i=0; i<numPixel; i++ )  
{  
    const uint16 offset = magicLUT [i] + time;  
    buffer[i] = texture[ offset & 0xffff ];  
}
```

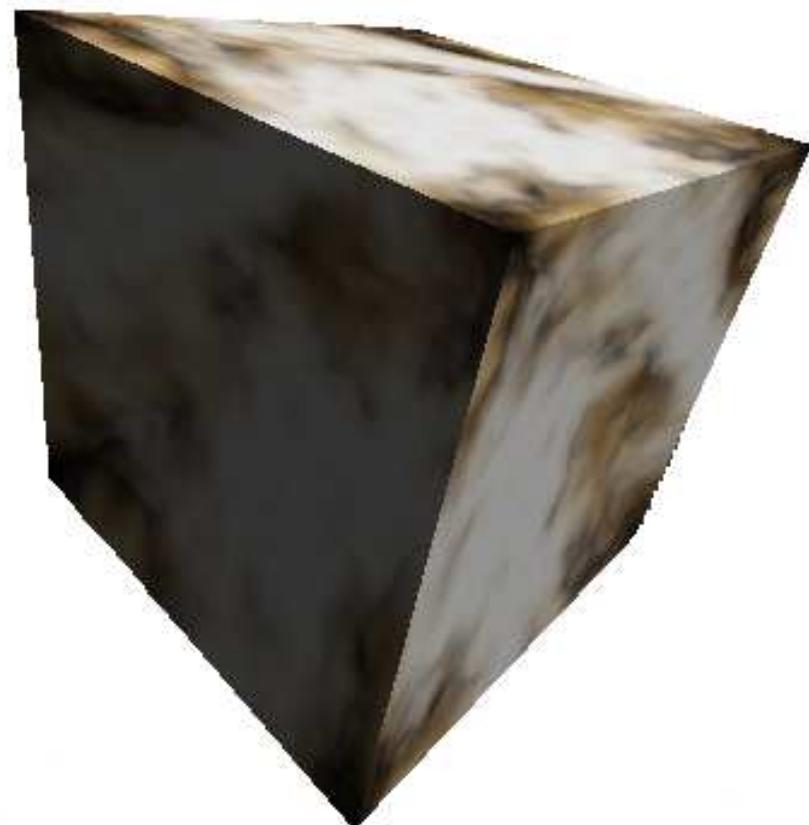
Pixel shader version

```
void main( void ) //for( each pixel p )  
{  
    vec2 offset = magicFormula(p, time);  
    gl_FragColor = texture2D(texture, offset);  
}
```



Old-school effects are back

- Others?
 - Rasterizers!
 - Vertex transformation
 - Triangle rasterization
 - Not perspective corrected
 - Metaballs
 - Plasmas
 - Raytracing



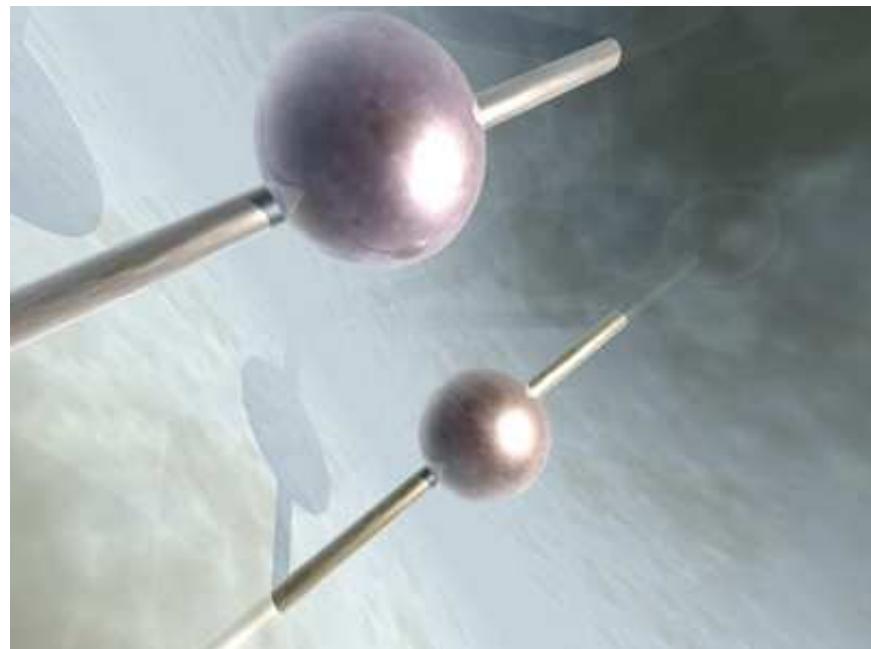


Old-school effects are back

- Whitted raytracing of simple scenes/primitives
 - A classic in demoscene
 - With fake analytic Ambient Occlusion



Chocolux, by Auld, 2007 [1 kbyte demo]



Kinderpainter, by rgba, at BCN 2006 [4k kbytes demo]

Old-school effects are back

```
#include <windows.h>
#include <GL/gl.h>
#include <GL/glext.h>

char *vsh="\
varying vec3 s[4];\
void main()\{\
gl_Position=gl_Vertex;\n
s[0]=vec3(0);\n
s[3]=vec3(sin(abs(gl_Vertex.x*.0001)),\n
cos(abs(gl_Vertex.x*.0001)),0);\n
s[1]=s[3].zxy;\n
s[2]=s[3].zzx;}\";

char *fsh="\
varying vec3 s[4];\
void main()\{\
float t,b,c,h=0;\
vec3 m,n,p=vec3(.2),d=normalize(.001*gl_FragCoord.rgb-p);\n
for(int i=0;i<4;i++){\n
t=2;\n
for(int i=0;i<4;i++){\n
b=dot(d,n=s[i]-p);\n
c=b*b+.2-dot(n,n);\n
if(b-c<t)if(c>0){m=s[i];t=b-c;}\n
}\n
p+=t*d;\n
d=reflect(d,n=normalize(p-m));\n
h+=pow(n.x*n.x,.44.)+n.x*n.x*.2;\n
}\n
gl_FragColor=vec4(h,h*h,h*h*h,h);}\";
```

```
PIXELFORMATDESCRIPTOR pfd={0,1,PFD_SUPPORT_OPENGL|PFD_DOUBLEBUFFER, 32, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 32, 0, 0, 0, 0, 0, 0, 0, 0};
DEVMODE dmScreenSettings={ 0,0,0,sizeof(DEVMODE),0,DM_PELSWIDTH|DM_PELSHEIGHT,
0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,1024,768,0,0,0,0,0,0,0,0,0,0,0,0,0,0};

void WinMainCRTStartup()
{
    ChangeDisplaySettings(&dmScreenSettings,CDS_FULLSCREEN);
    HDC hDC = GetDC(CreateWindow("edit",0,WS_POPUP | WS_VISIBLE | WS_MAXIMIZE, 0,
0, 0, 0, 0, 0, 0));
    SetPixelFormat(hDC, ChoosePixelFormat(hDC, &pfd), &pfd);
    wglMakeCurrent(hDC, wglCreateContext(hDC));
    ShowCursor(0);
    GLuint p = ((PFNGLCREATEPROGRAMPROC)wglGetProcAddress("glCreateProgram"))();
    GLuint s = ((PFNGLCRAE_SHADERPROC)
wglGetProcAddress("glCreateShader"))(GL_VERTEX_SHADER);
    ((PFNGLSHADERSOURCEPROC)wglGetProcAddress("glShaderSource"))(s,1, &vsh,0);
    ((PFNGLCOMPILESHADERPROC)wglGetProcAddress("glCompileShader"))(s);
    ((PFNGLATTACHSHADERPROC)wglGetProcAddress("glAttachShader"))(p,s);
    s = ((PFNGLCRAE_SHADERPROC)
wglGetProcAddress("glCreateShader"))(GL_FRAGMENT_SHADER);
    ((PFNGLSHADERSOURCEPROC)wglGetProcAddress("glShaderSource"))(s,1, &fsh,0);
    ((PFNGLCOMPILESHADERPROC)wglGetProcAddress("glCompileShader"))(s);
    ((PFNGLATTACHSHADERPROC)wglGetProcAddress("glAttachShader"))(p,s);
    ((PFNGLLINKPROGRAMPROC)wglGetProcAddress("glLinkProgram"))(p);
    ((PFNGLUSEPROGRAMPROC)wglGetProcAddress("glUseProgram"))(p);
loop:
    int t=GetTickCount();
    glRecti(t,t,-t,-t);
    SwapBuffers(hDC);
    if (GetAsyncKeyState(VK_ESCAPE)) ExitProcess(0);
    goto loop;
}
```

Source code of chocolux, by Auld (link with Crinkler)

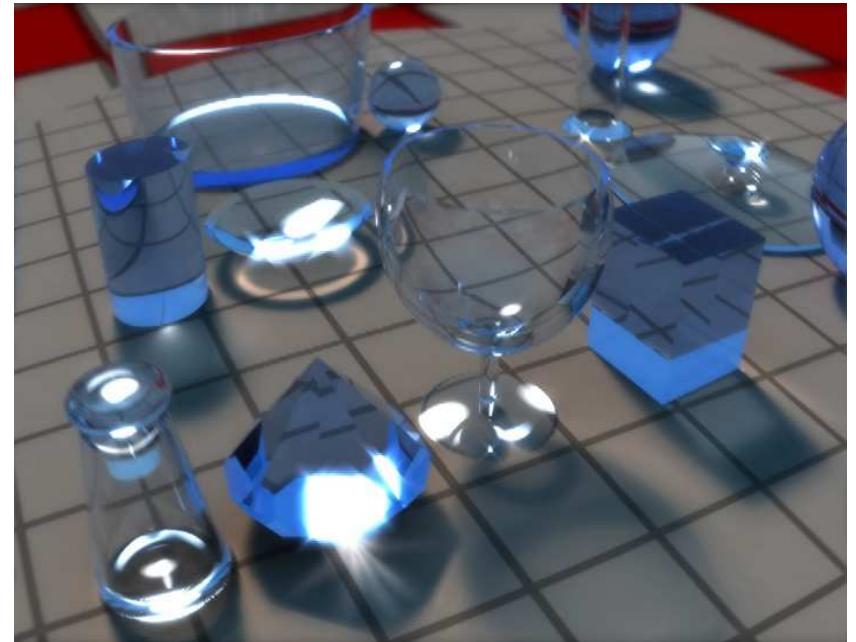


Old-school effects are back

- Pathtracing of simple scenes/primitives



Off the shelf, by Loonies, at Breakpoint 2008 [4k kbytes image]



PhotonRace, by Archee, at Buentzli 2008 [4k kbytes image]



Old-school effects are back

- GPU raytracing beyond spheres and planes (I mean polygons)



Images reproduced with permission of Vrcontext (www.vrcontext.com)



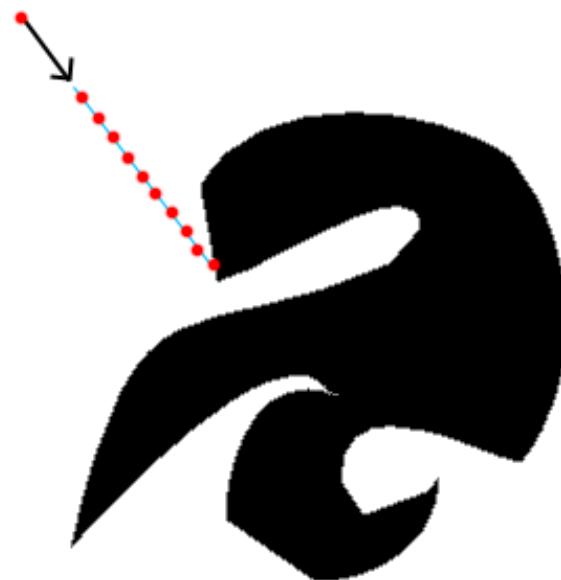
Old-school effects are back

- GPU raytracing beyond spheres and planes (I mean polygons)
 - A very hot research topic today (because raytracing is the future...)
 - Difficult to beat CPU raytracers
 - kd-tree/bih/bvh traversal is quite incoherent
 - They all need a stack (unavailable today on shaders).
 - For massive models, streaming to video memory is needed. That makes it more complex.
 - In any case, demosceners have not been interested on real raytracing so far; even less in the 4k categories.



Old-school effects are back

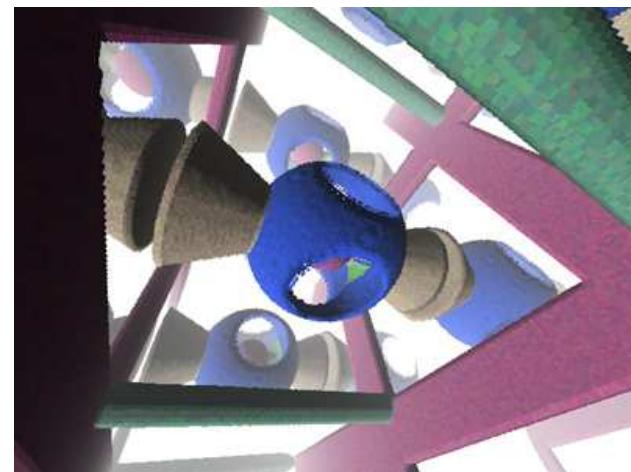
- Raymarching
 - Kind of raytracing for all those objects that don't have an analytic intersection function.





Old-school effects are back

- Raymarching -- what?
 - Heightmaps
 - Volume textures
 - Procedural isosurfaces
 - Analytic surfaces



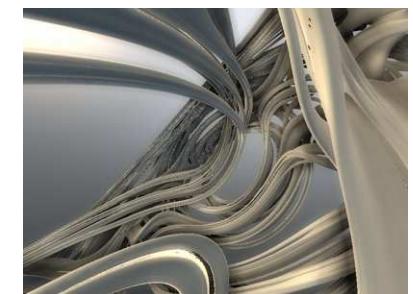
3D texture volume raymarching (*rgba*)



Heightmap raymarching. Hymalaya, by TBC 2008, 1 kbytes demo



Procedural isosurface

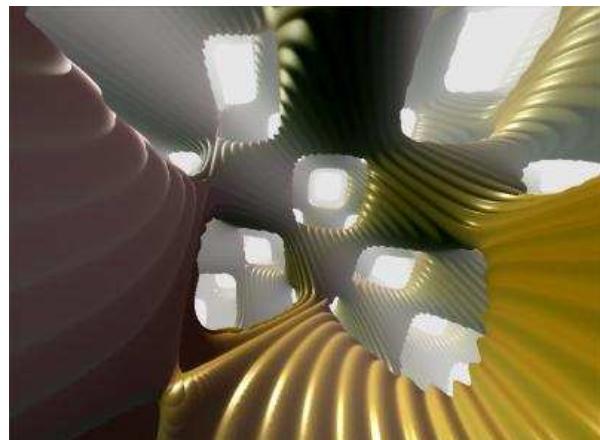


Analytic surface

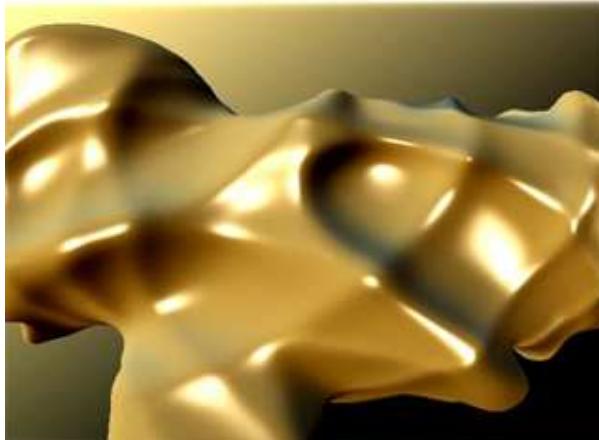


Old-school effects are back

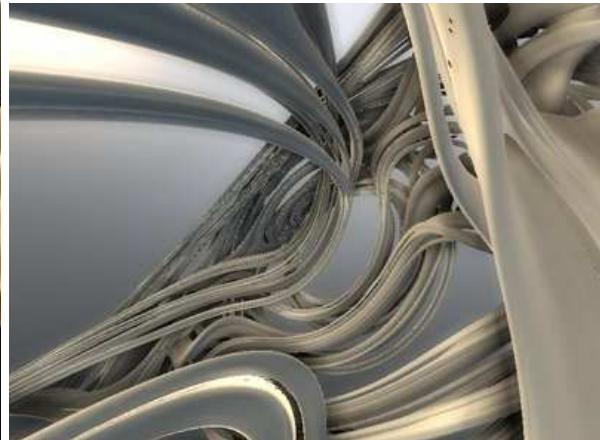
- Raymarching -- how?
 - Constant steps
 - Root finders (bisection, Newton-Raphson...)
 - Distance fields



Fality, by Loonies, 2006, a 4 kbytes demo



Tracie, by TBC, 2007, a 1 kbytes demo



Kindernoiser, by rgba, 2007, a 4 kbytes demo



Index

- Old school effects are back
- Rendering with distance fields



Rendering with distance fields

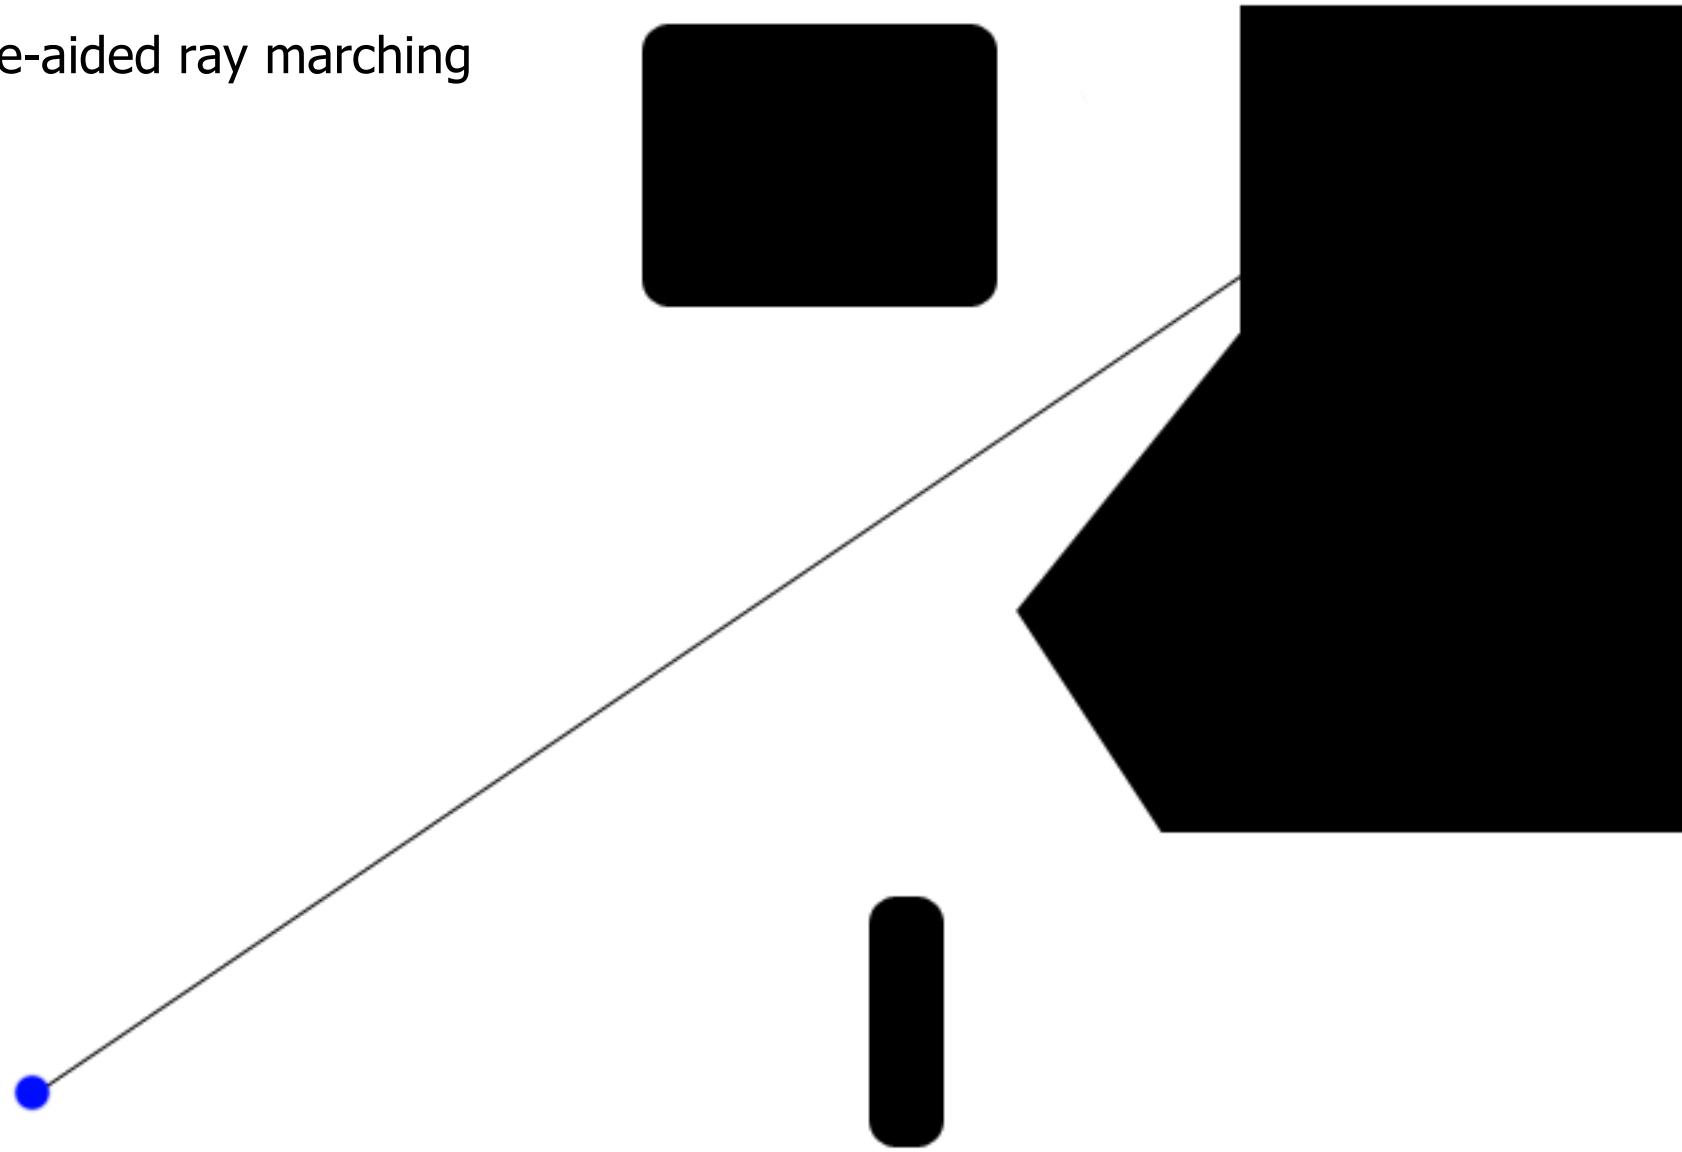


Slisesix, by rgba, at EuskalEncounter 2008 [4k kbytes image]

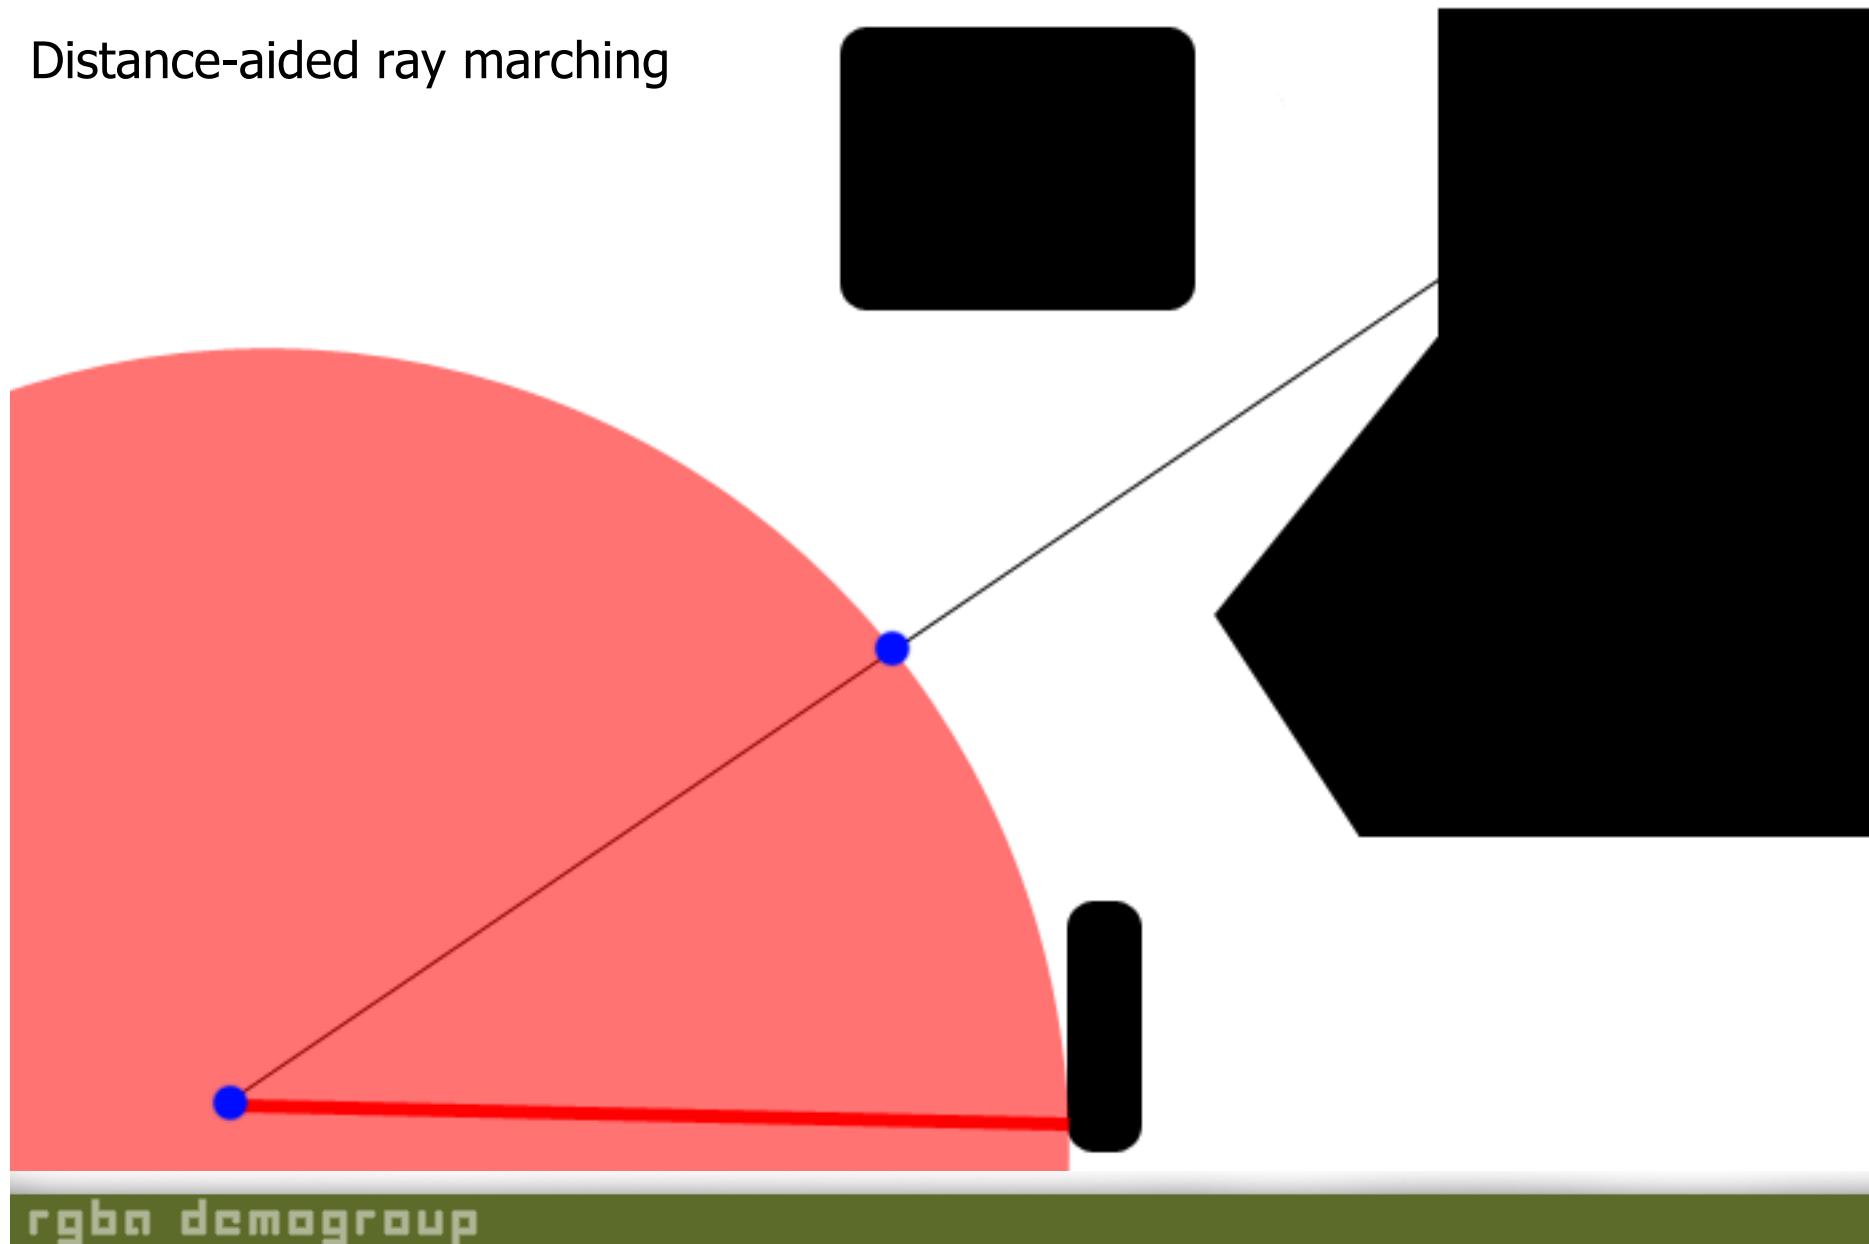
Rendering with distance fields

- Similarly previous works
 - “Ray tracing deterministic 3-D fractals” published at Siggraph 1989 by D.J.Sandin and others.
 - “Per-pixel displacement mapping with distance functions”, appeared in GPU Gems 2 (2005) by W.Donnelly.
- The trick is to be able to compute or estimate (a lower bound of) the distance to the closest surface at any point in space.
- This allows for marching in large steps along the ray.

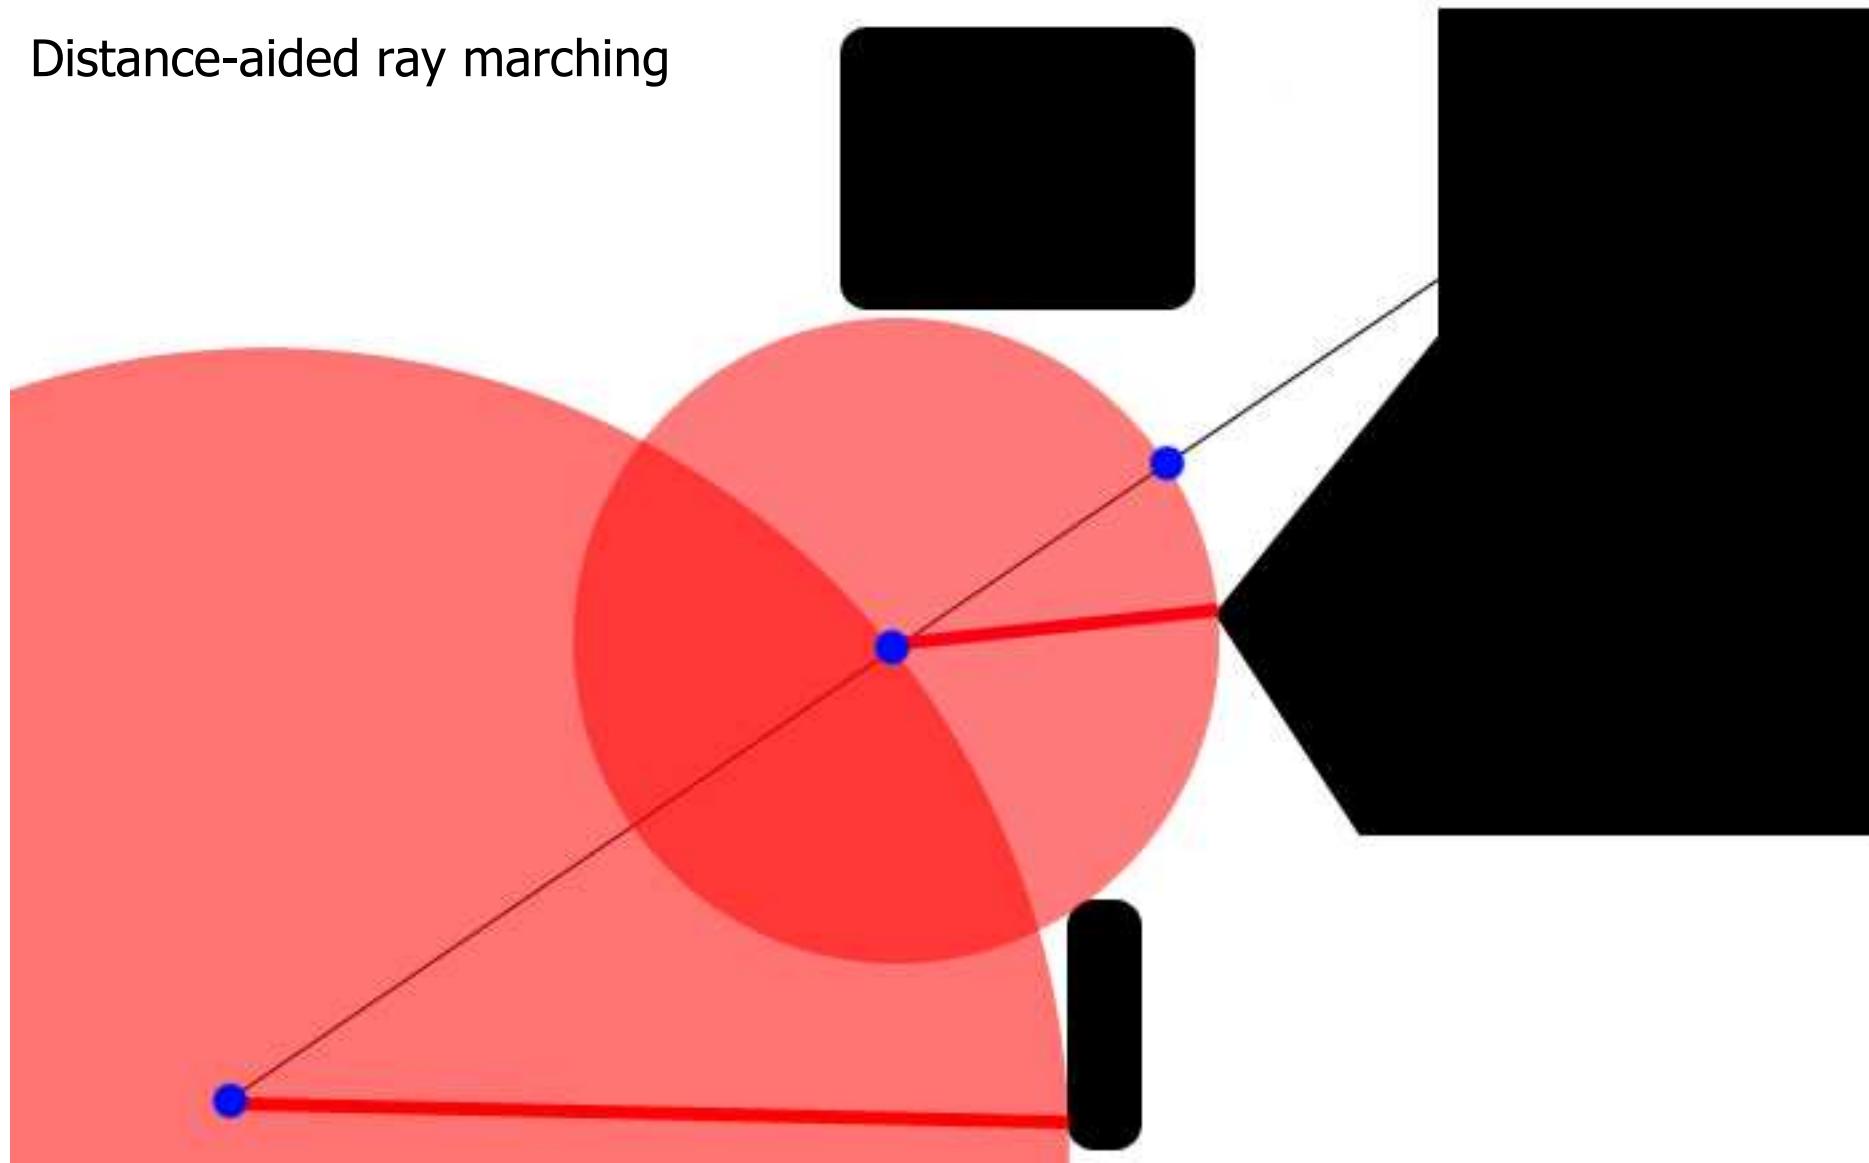
Distance-aided ray marching



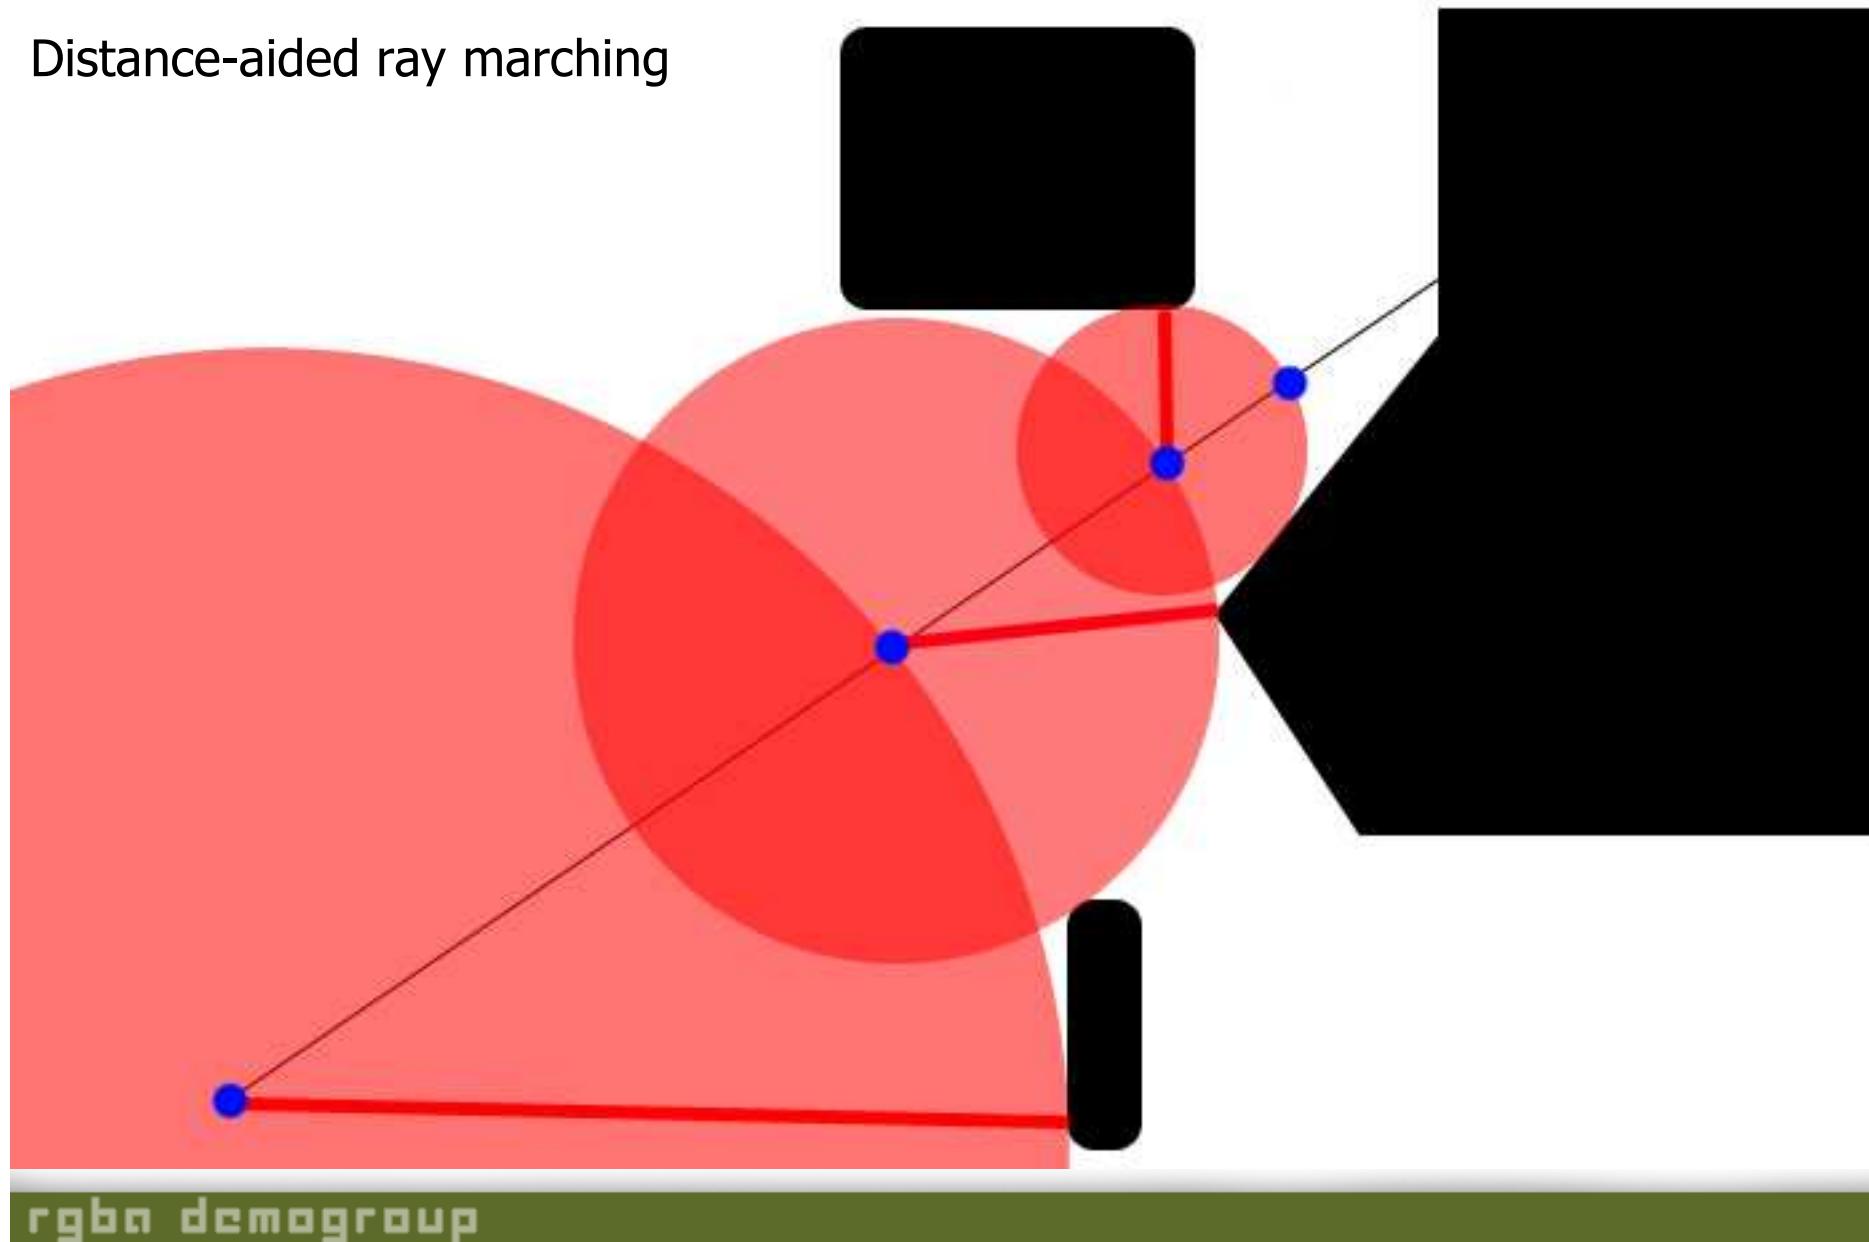
Distance-aided ray marching



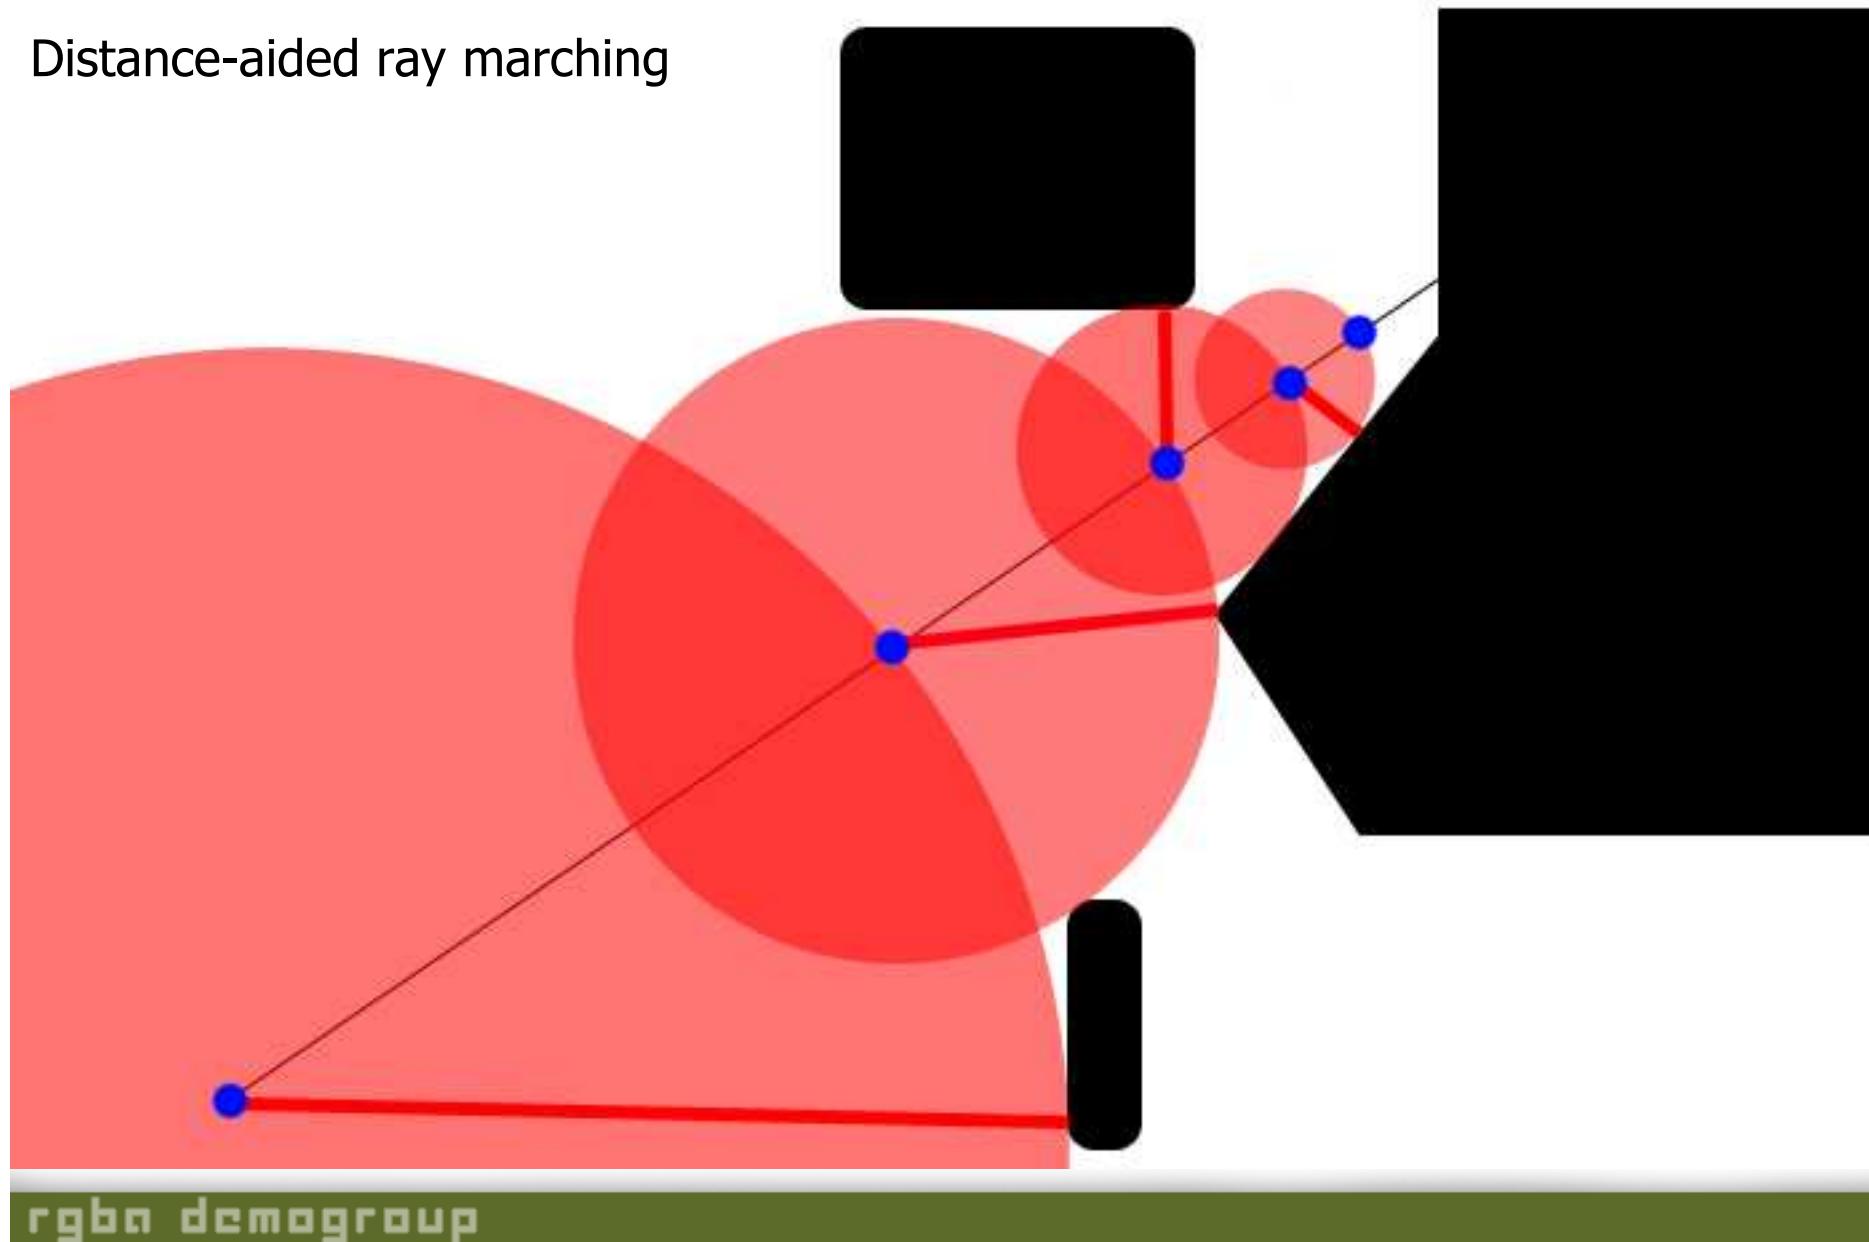
Distance-aided ray marching



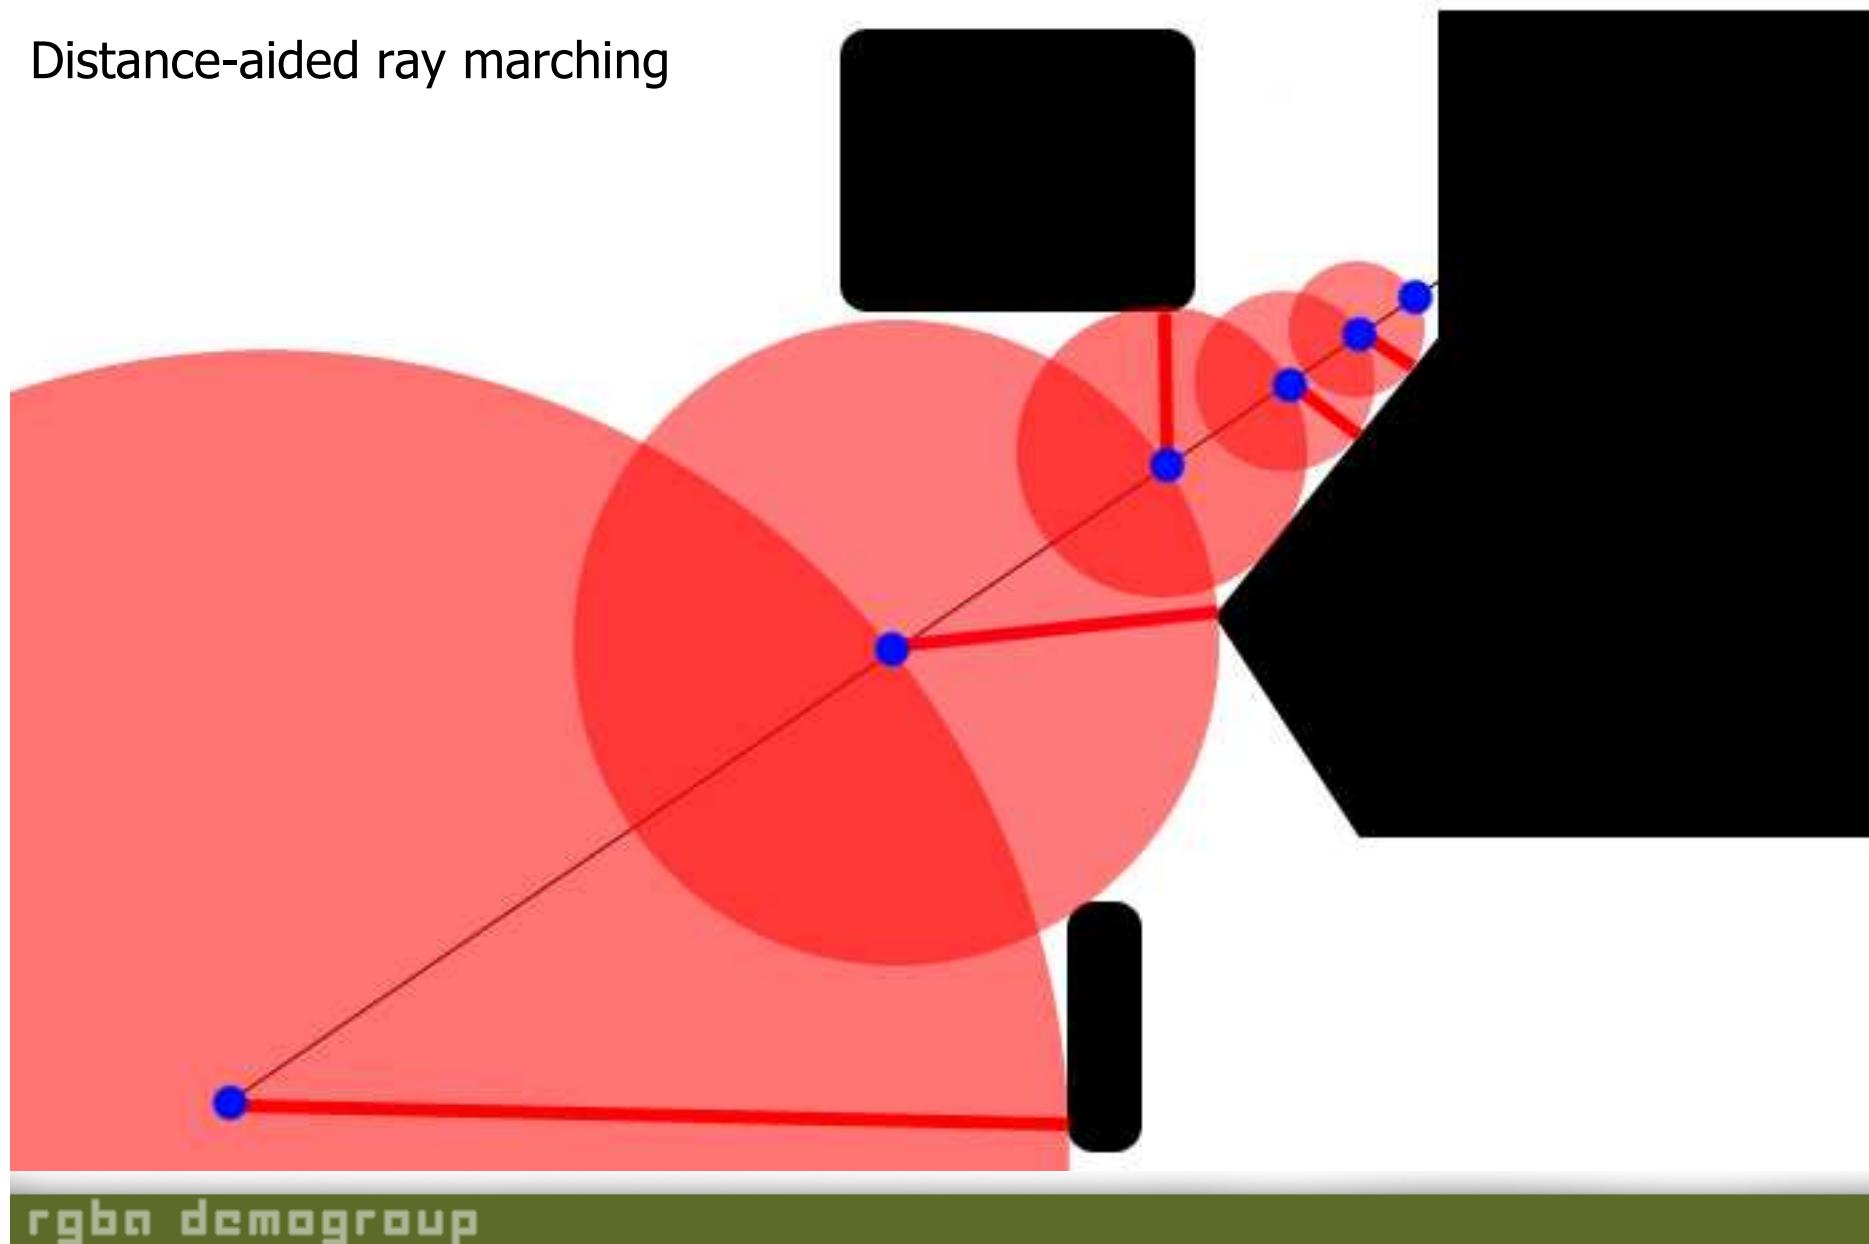
Distance-aided ray marching



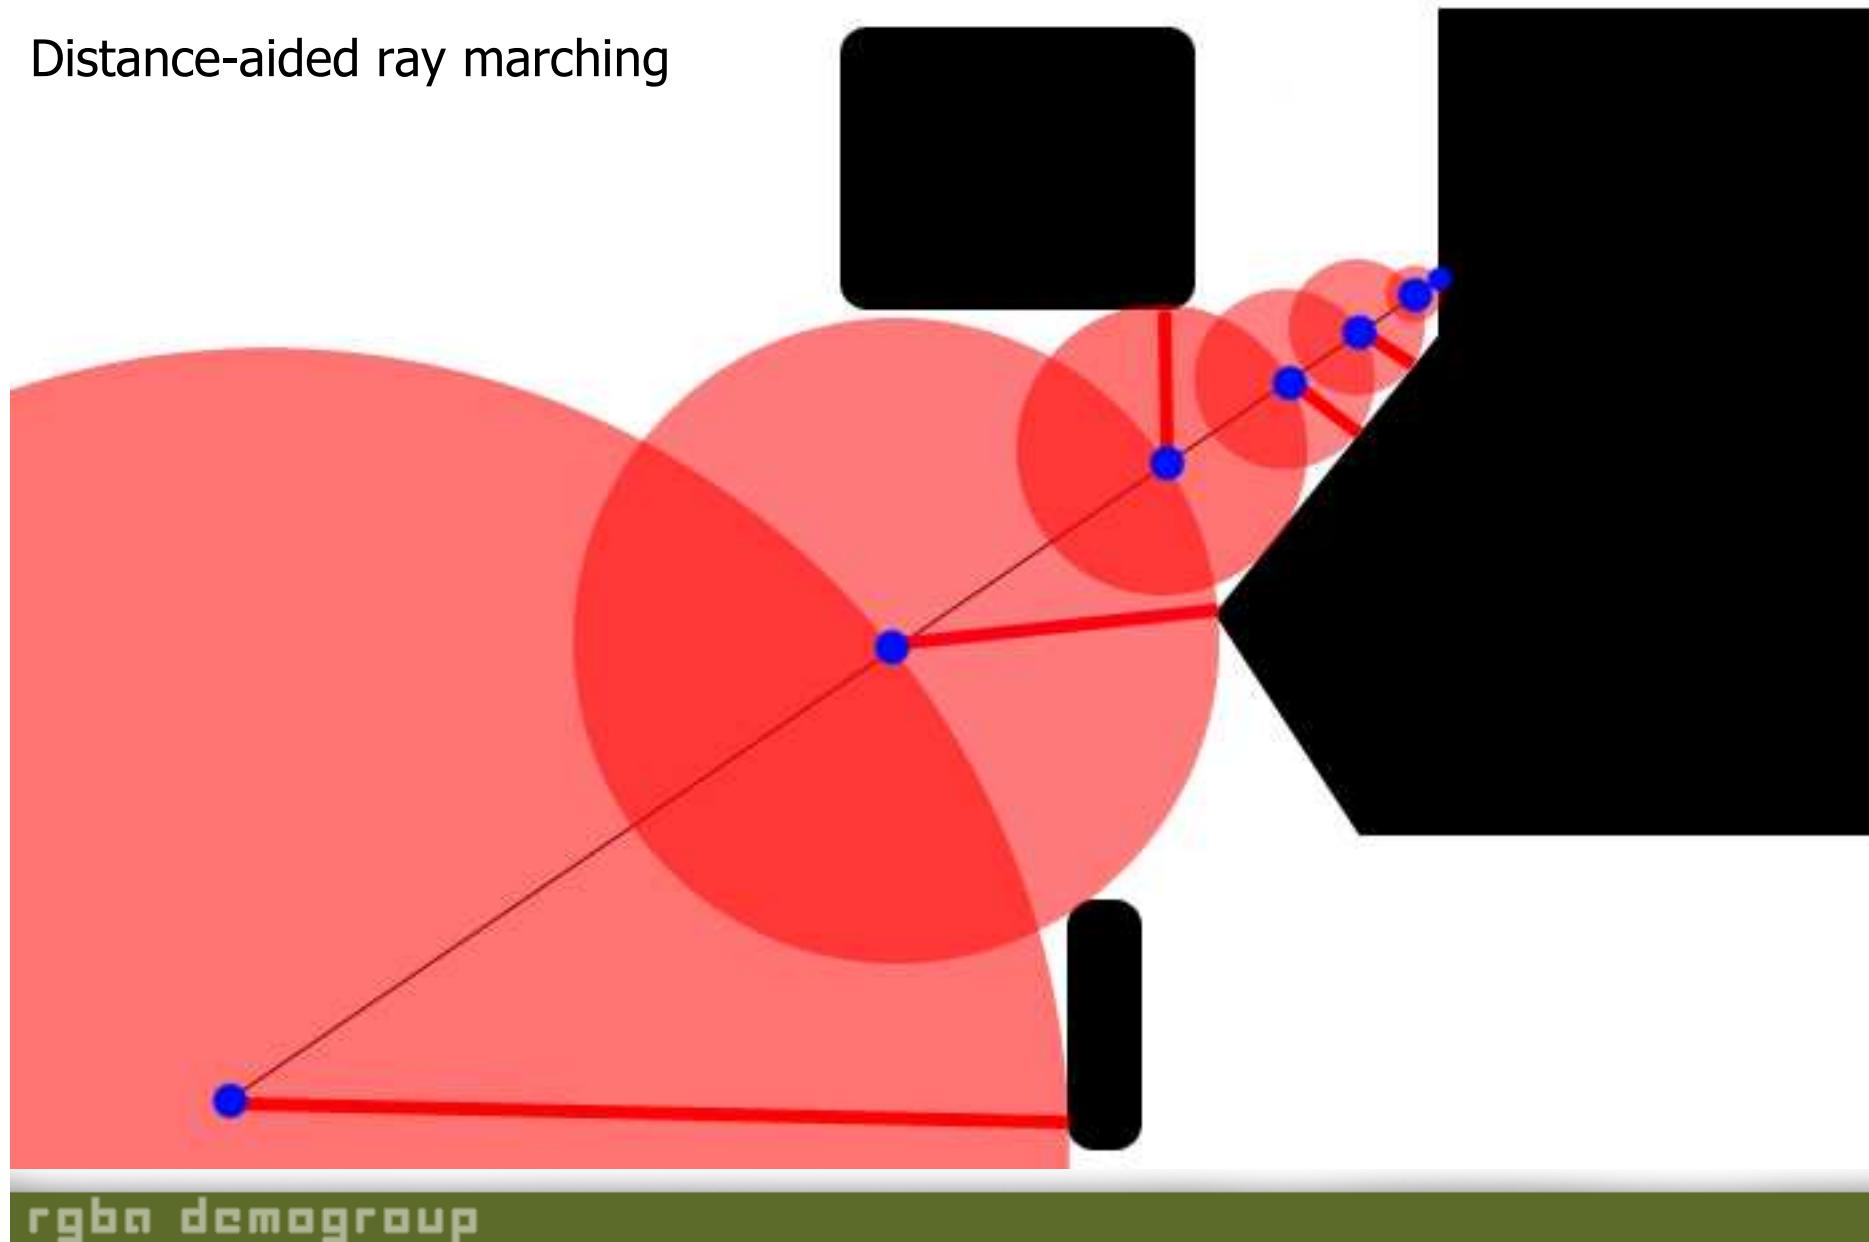
Distance-aided ray marching



Distance-aided ray marching



Distance-aided ray marching





Rendering with distance fields

Pros

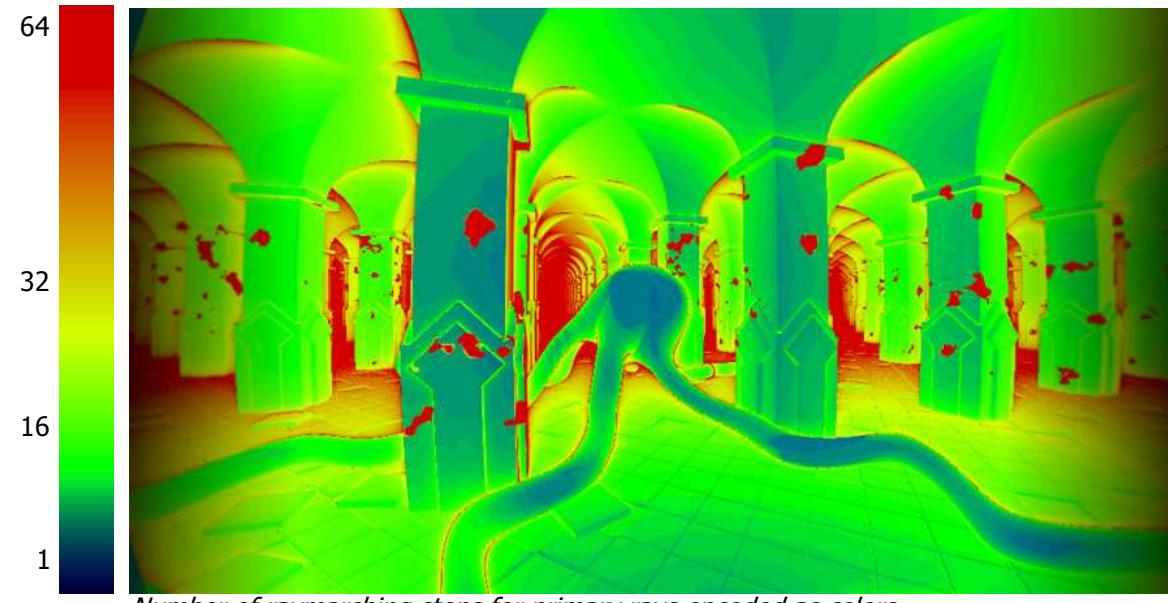
- Much faster than constant-size stepping.
- Much easier to control than root finders (bisection, Newton...)
- Room for optimization, like using bigger steps when we are further from the ray origin
 - Error in world coordinates decreases as $1/d$
 - So stepping proportionally to d results in constant screen space error.

Cons

- Slow on the boundaries of the objects (hopefully not that many pixels).
- Can control it by imposing a minimum step size.

Rendering with distance fields

- Slisesix needs 50 million evaluations of the very expensive distance function for a 1280x720 pixel image.
- 60% of the evaluations are for primary rays (av. 17 steps per ray).
- 40% of the evaluations are for lighting and shading.
- Note the very expensive marching on the object edges.





Rendering with distance fields

- We need a distance field:
 - Analytic computation ("Ray tracing deterministic 3-D fractals")
 - Precomputed (static scene) LUT
 - 3D texture ("Per pixel displacement mapping with distance functions")
 - Octree / KdTree
 - What if we do it 100% procedurally? ("Slisesix")



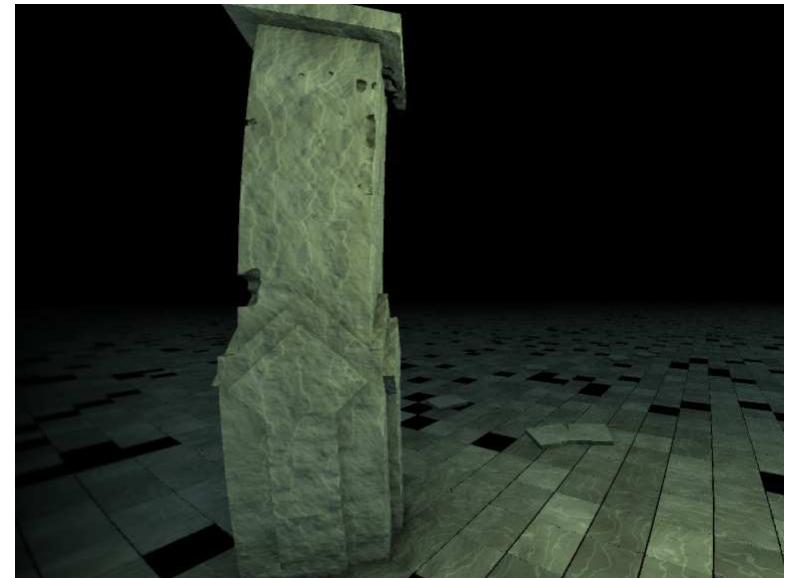
Rendering with distance fields

- Procedural distance fields
 - Don't define the surface first and then compute the distance field, but directly code a distance field and a surface will emerge.
 - Tweak the distance field directly until you get what you want/can.
 - Helpful techniques that can be used:
 - Arbitrary combination and instantiation
 - Inifinite repetition
 - Deforming space: twisting, bending, deforming
 - Cheap detail surfaces
 - Blend shapes

Rendering with distance fields :: Combination

- Combination of (instanced) distance fields can be done by taking the min of the distance fields involved.
- Instance transformation can be done by inverse transforming the domain (the input to the distance function).

```
float combinedDistanceField( vec3 p )
{
    float dist1 = distanceField_A( M1inv*p );
    float dist2 = distanceField_A( M2inv*p );
    float dist3 = distanceField_B( M3inv*p );
    return min( dist1, min( dist2, dist3 ) );
}
```





Rendering with distance fields :: Domain repetition

- dist = fourMagicColumns(p.x, p.y, p.z);





Rendering with distance fields :: Domain repetition

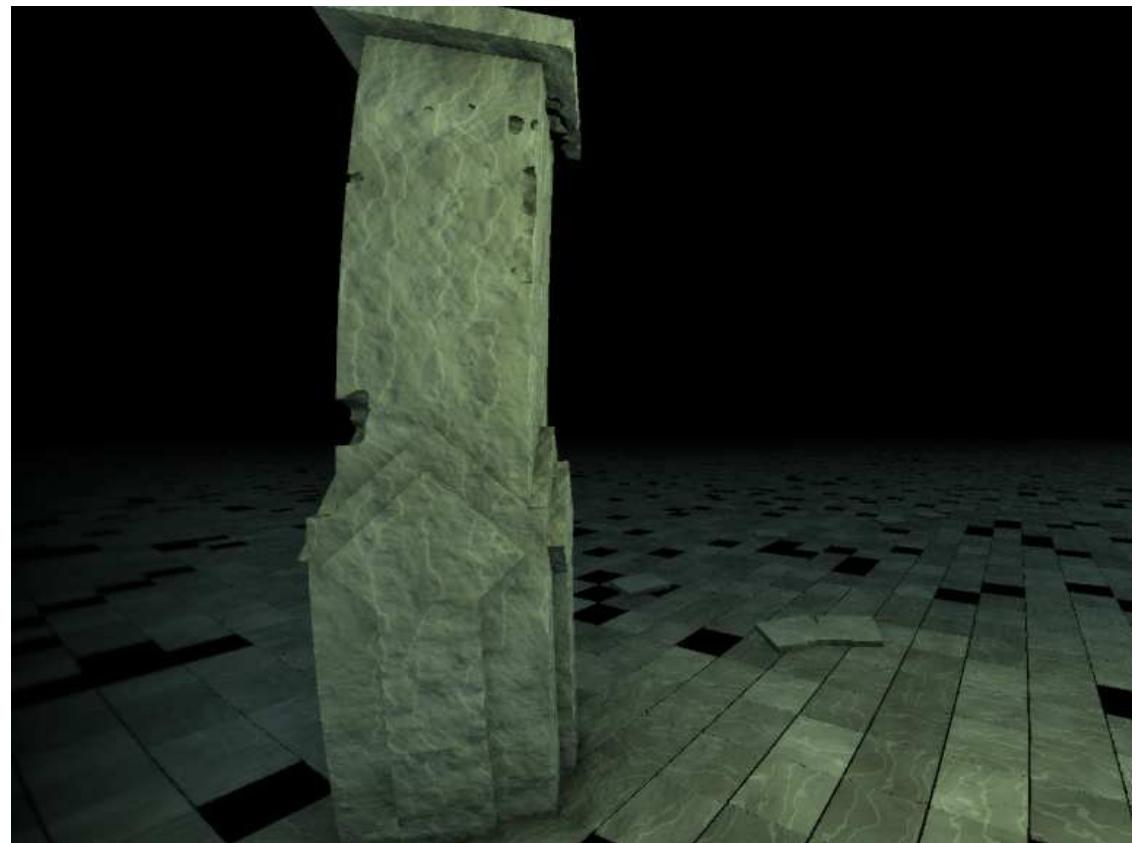
- dist = fourMagicColumns(mod(p.x,1), p.y, mod(p.z,1));





Rendering with distance fields :: Domain distortion

```
float dist = distanceToColumn(p);
```





Rendering with distance fields :: Domain distortion

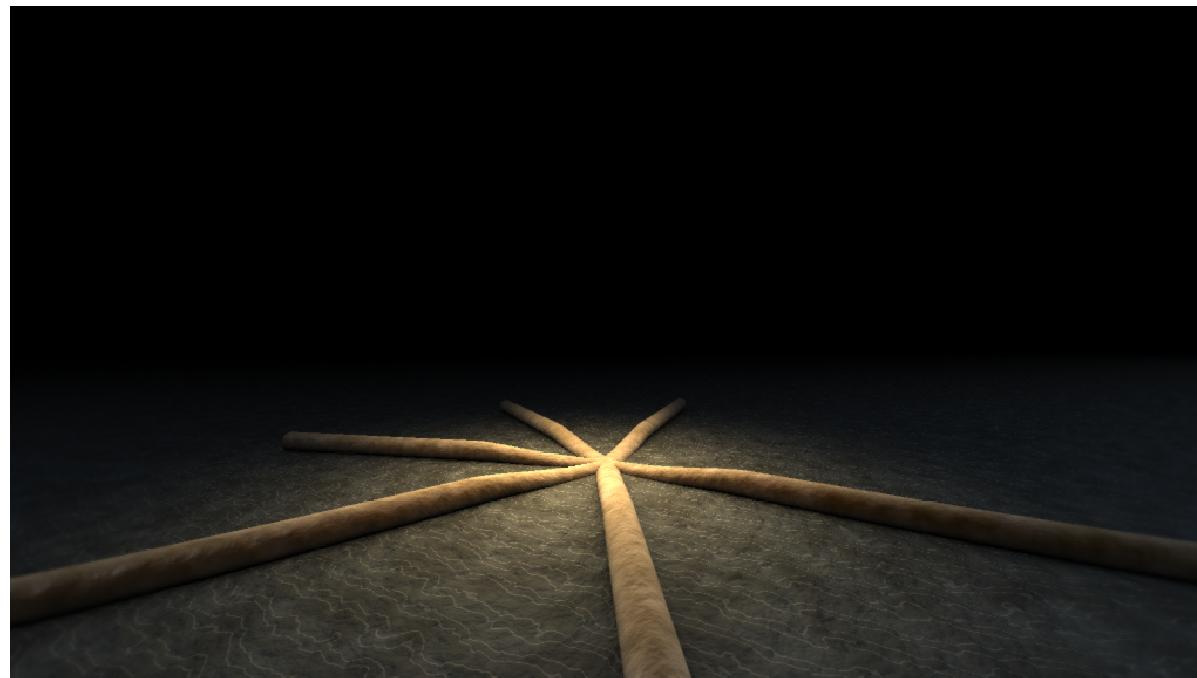
```
float twistedColumn( vec3 p )
{
    vec3 q = rotateY(p, p.y*1.7);
    return distanceToColumn(q);
}
```





Rendering with distance fields :: Domain distortion

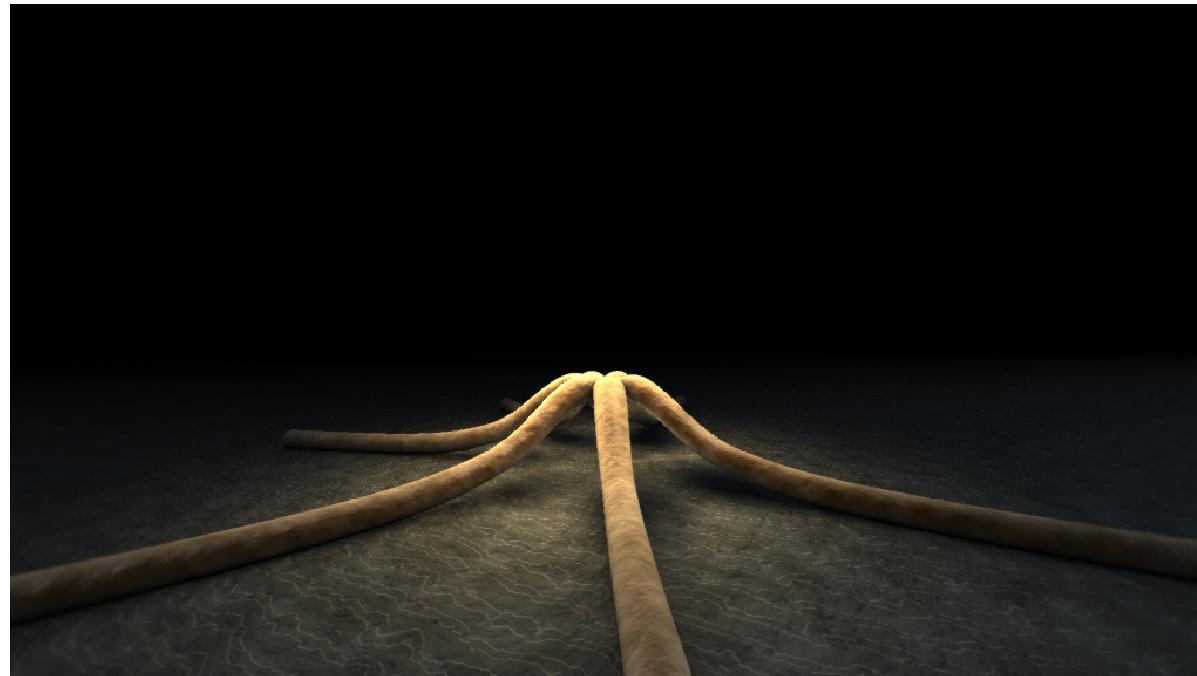
```
float rr = dot(p.xy,p.xy);
for( int i=0; i<6; i++ )
{
    vec3 q = rotateY( p, TWOPI*i/6.0 );
    distance = min( distance, distanceToTheXAxis(q) );
}
```





Rendering with distance fields :: Domain distortion

```
float rr = dot(p.xy,p.xy);
for( int i=0; i<6; i++ )
{
    vec3 q = rotateY( p, TWOPI*i/6.0 );
    q.y += 0.6*rr*exp2(-10.0*rr);
    distance = min( distance, distanceToTheXAxis(q) );
}
```



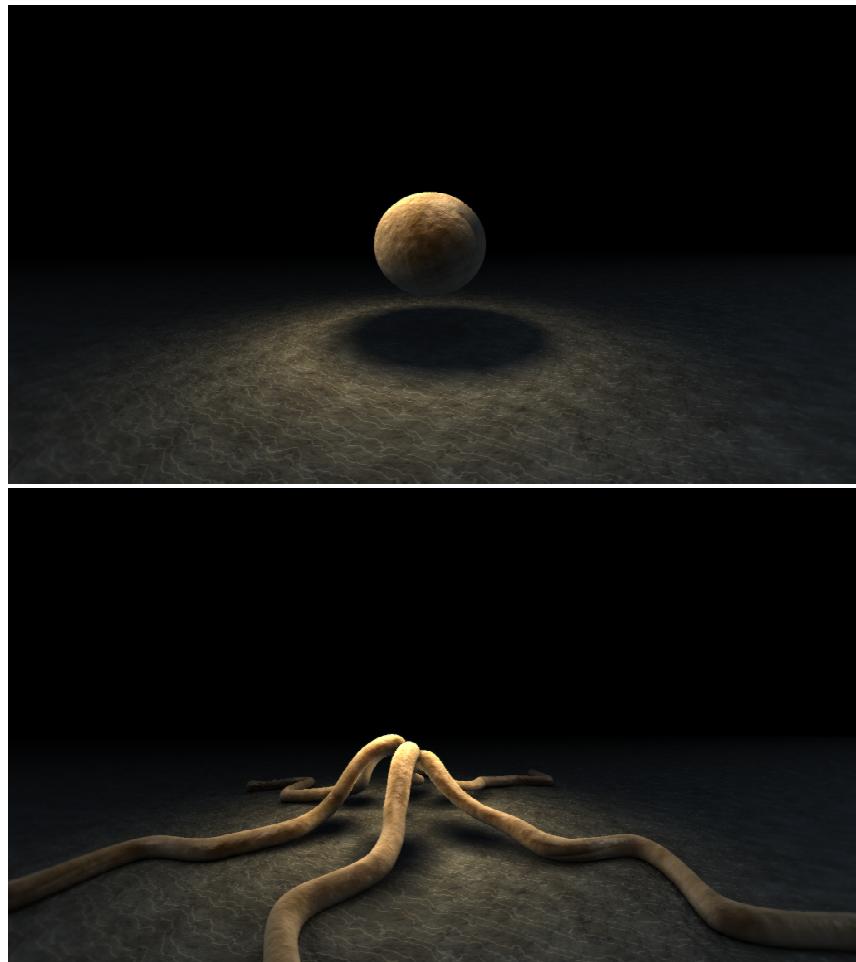


Rendering with distance fields :: Domain distortion

```
float rr = dot(p.xy,p.xy);
for( int i=0; i<6; i++ )
{
    vec3 q = rotateY( p, TWOPI*i/6.0 + 0.4*rr*noise2f(vec3(4*rr,6.3*i)) );
    q.y += 0.6*rr*exp2(-10.0*rr);
    distance = min( distance, distanceToTheXAxis(q) );
}
```



Rendering with distance fields :: Blending fields



```
float distanceToMonster( vec3 p )  
{  
    float dist1 = distanceToBall(p);  
    float dist2 = distanceToTentacles(p);  
    float bfact = smoothstep( length(p), 0, 1 );  
    return mix( dist1, dist2, bfact );  
}
```



Rendering with distance fields :: Adding details

```
dist = distanceToColmuns(p);
```





Rendering with distance fields :: Adding details

```
dist = distanceToColmuns(p) + 0.000001*clamp(fbm(p), 0, 1);
```





Rendering with distance fields :: Lighting

- Lighting
 - Normals
 - Bump mapping
 - Soft shadows
 - Ambient Occlusion



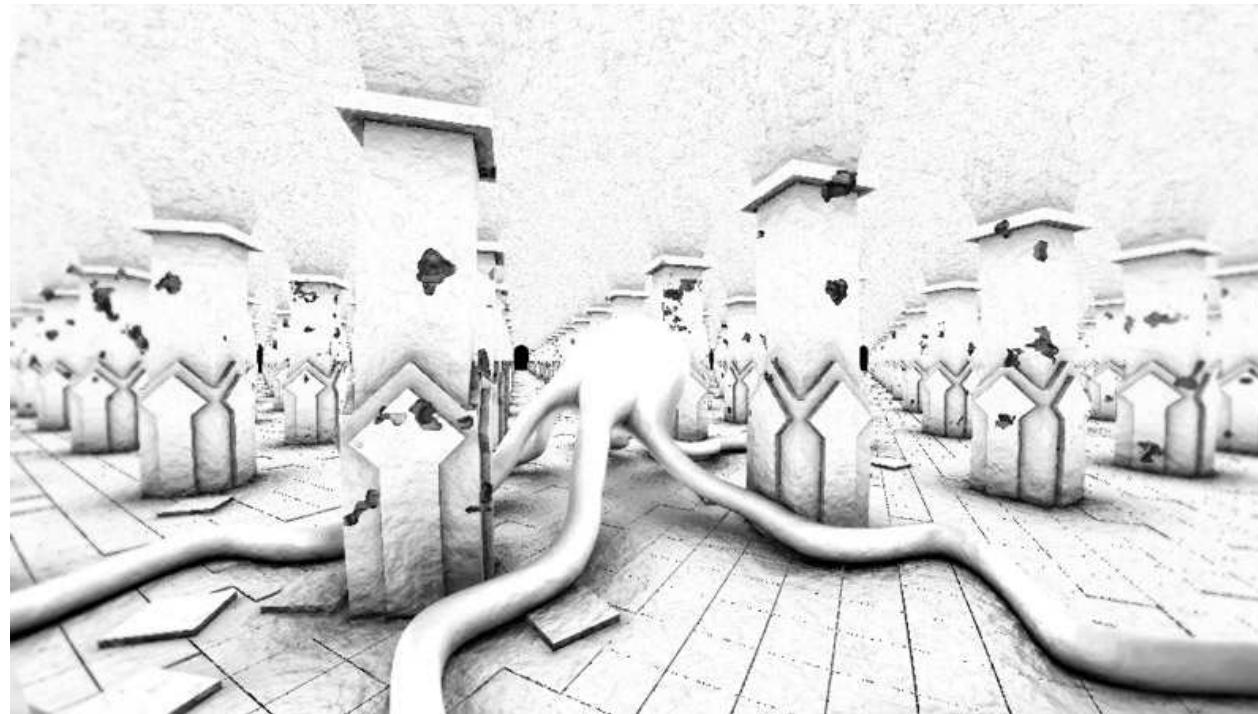
Rendering with distance fields :: Lighting

- Normals computed by central differences on the distance field at the shading point (gradient approximation).
- Bump map computed by adding the gradient of a fractal sum of Perlin noise functions to the surface normal.
 - $n = \text{normalize}(\text{grad}(\text{distance}, p) + \text{bump} * \text{grad}(\text{fbm}, p));$
 - *bump* is small and depend on the material.
 - $\text{grad}(\text{func}, p) = \text{normalize}($
 $\text{func}(p+\{\text{eps}, 0, 0\}) - \text{func}(p-\{\text{eps}, 0, 0\}),$
 $\text{func}(p+\{0, \text{eps}, 0\}) - \text{func}(p-\{0, \text{eps}, 0\}),$
 $\text{func}(p+\{0, 0, \text{eps}\}) - \text{func}(p-\{0, 0, \text{eps}\}));$



Rendering with distance fields :: Ambient Occlusion

- Fake and fast Ambient Occlusion.
- VERY CHEAP, even cheaper than primary rays! Only 5 distance evaluations instead of casting thousand of rays/evaluations.

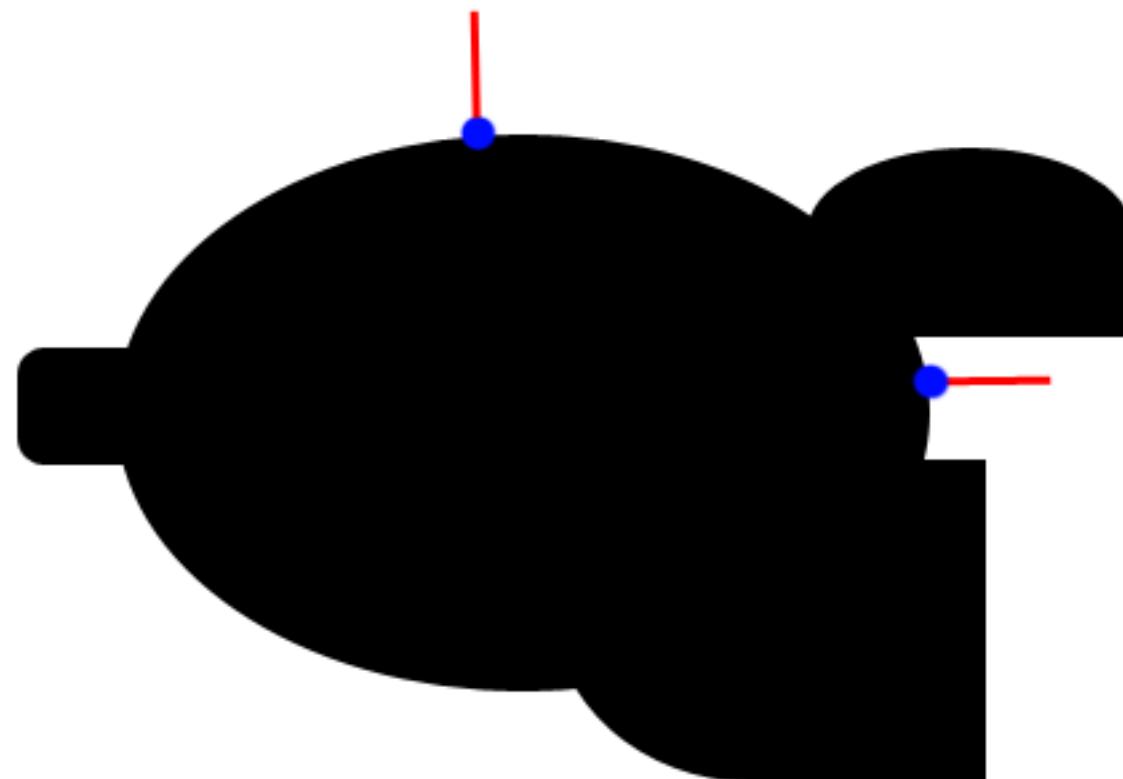




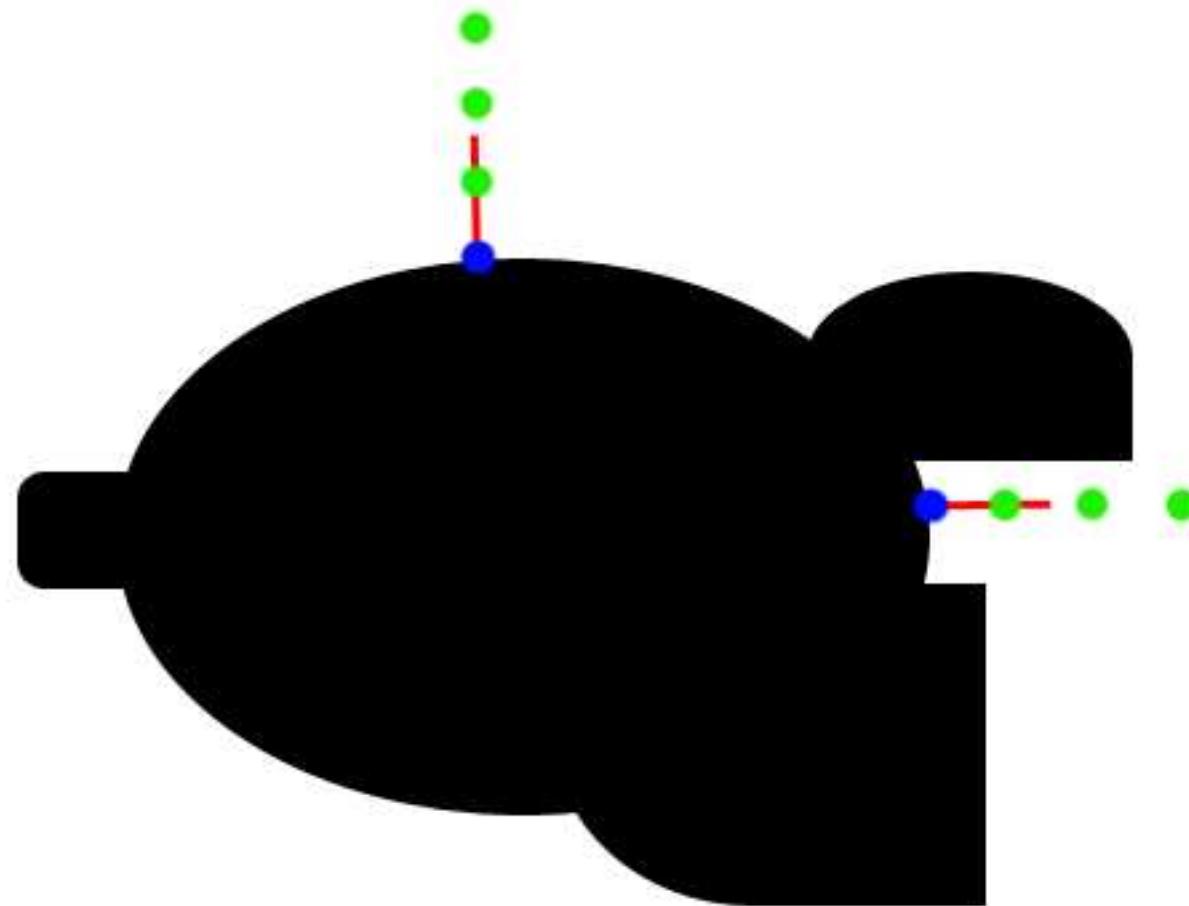
Rendering with distance fields :: Ambient Occlusion

- In a regular raytracer, primary rays/AO cost is 1:2000. Here, it's 3:1 (that's almost four orders of magnitude speedup!).
- It's NOT the screen space trick (SSAO), but 3D.
- The basic technique was invented by Alex Evans, aka Statix ("Fast Approximation for Global Illumination on Dynamic Scenes", 2006). Greets to him!
- The idea: let p be the point to shade. Sample the distance field at a few (5) points around p and compare the result to the actual distance to p . That gives surface proximity information that can easily be interpreted as an (ambient) occlusion factor.

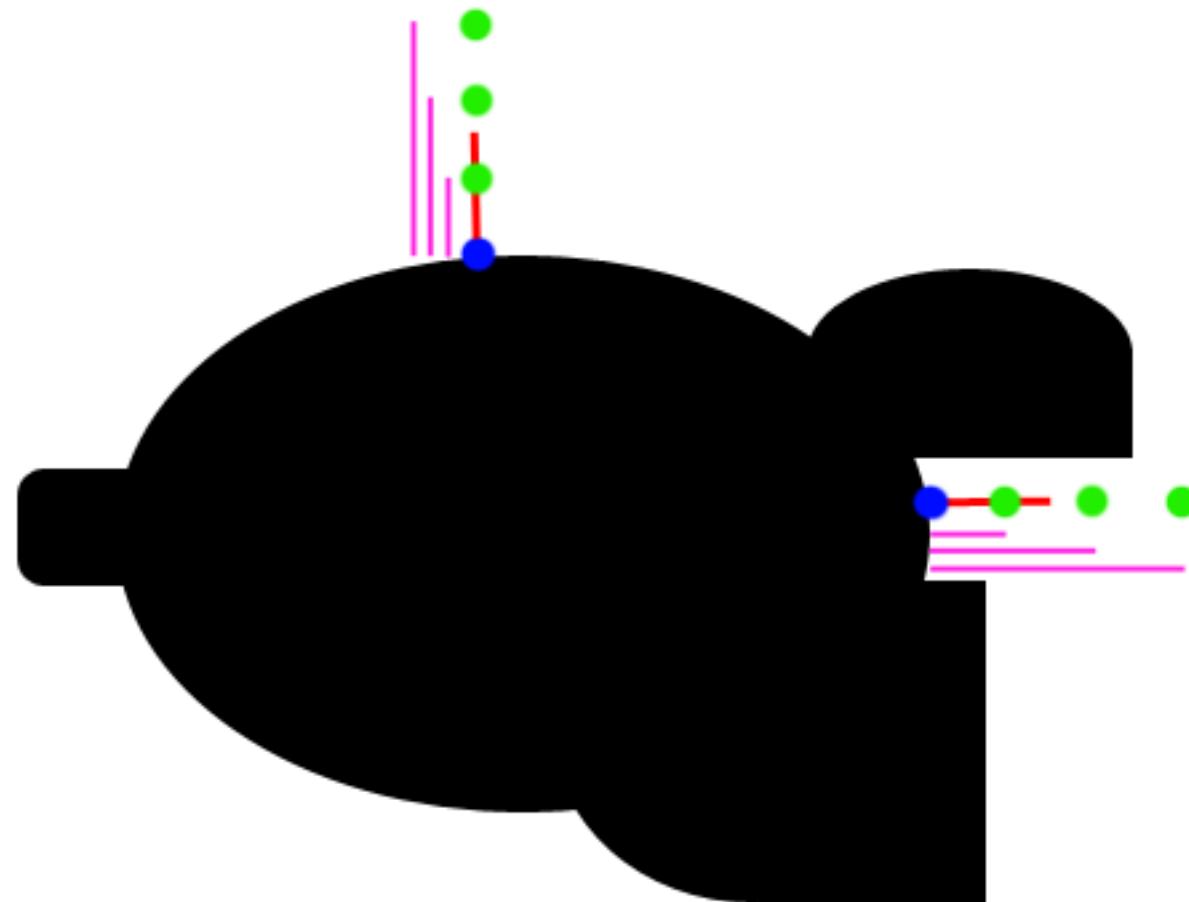
Rendering with distance fields :: Ambient Occlusion



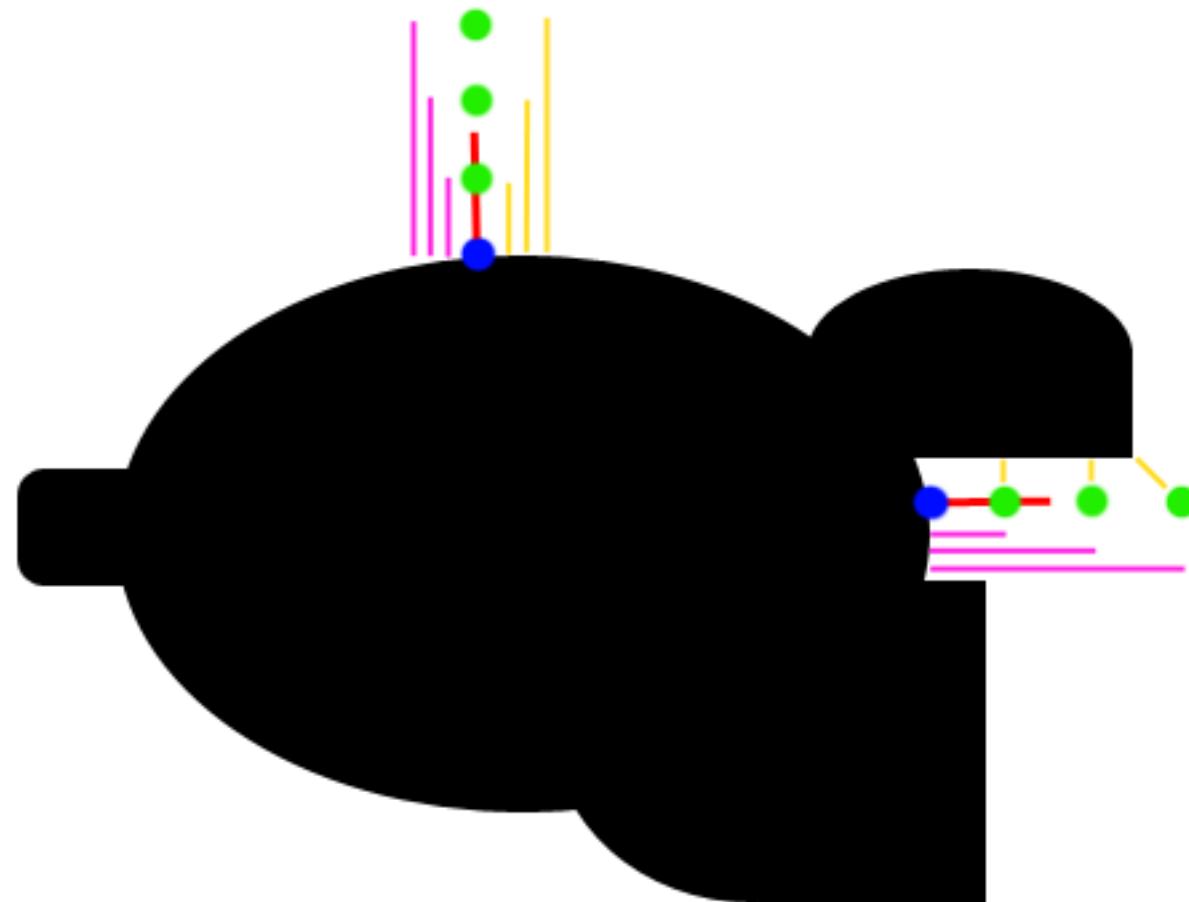
Rendering with distance fields :: Ambient Occlusion



Rendering with distance fields :: Ambient Occlusion



Rendering with distance fields :: Ambient Occlusion

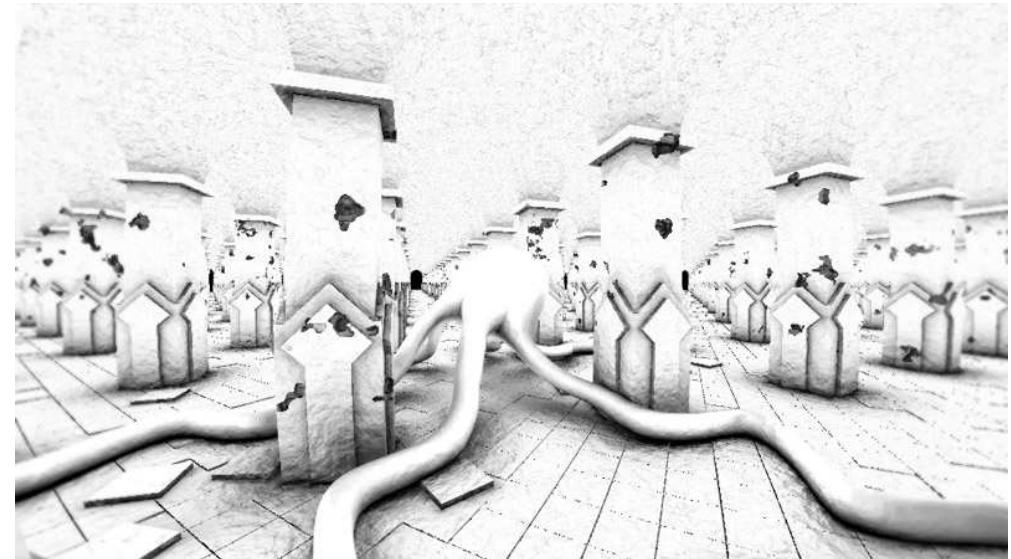
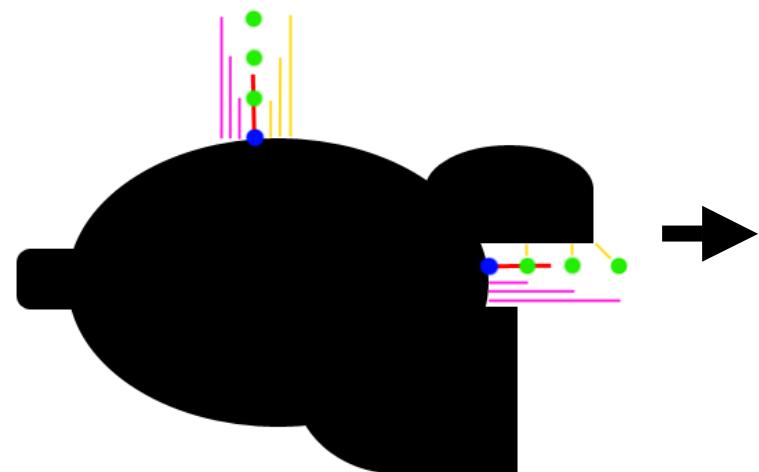


Rendering with distance fields :: Ambient Occlusion

$$ao = 1 - k \cdot \sum_{i=1}^5 \frac{1}{2^i} (pink_i - yellow_i)$$

$$ao = 1 - k \cdot \sum_{i=1}^5 \frac{1}{2^i} (i \cdot \Delta - distfield(p + n \cdot i \cdot \Delta))$$

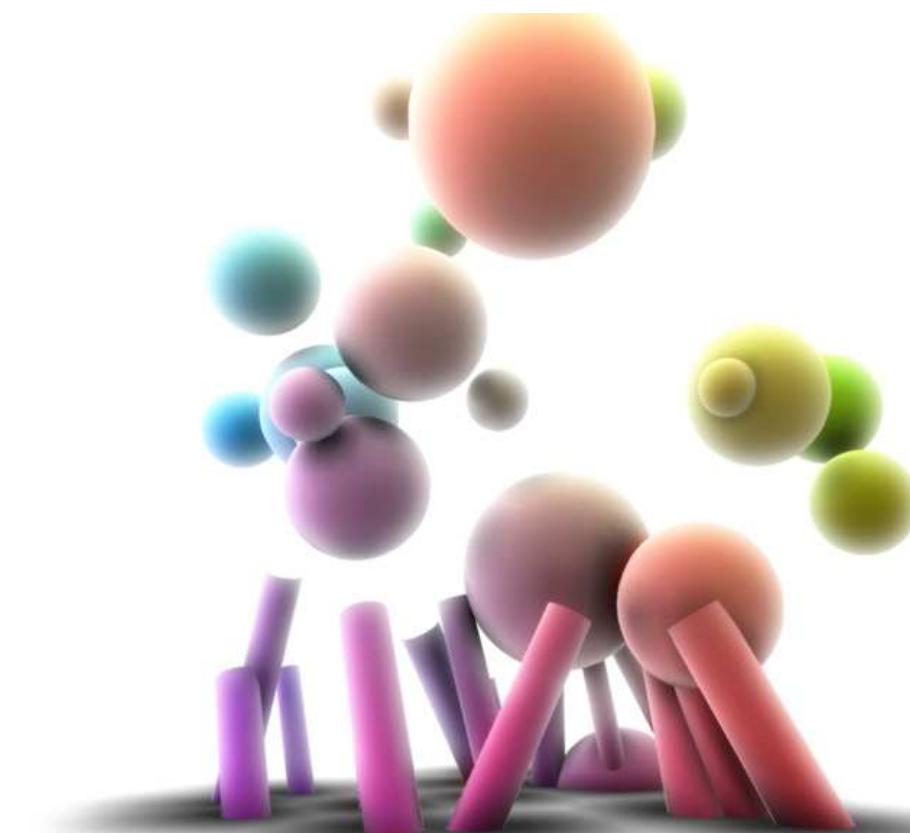
- The exponential decay is there so further away surfaces occlude less than near by ones.





Rendering with distance fields :: Ambient Occlusion

- Works in realtime too, provided you can compute distances to surfaces.





Rendering with distance fields :: Soft Shadows

- Fake and fast soft shadows.
- Only 6 distance evaluations used instead of casting hundreds of rays.
- Pure geometry-based, not bluring.
- Recipe: take n points on the line from the surface to the light and evaluate the distance to the closest geometry. Find a magic formula to blend the n distances to obtain a shadow factor.





Rendering with distance fields

- On a GeForce 8800 GTX, it renders around 20 times faster than on a dual core CPU. It will very soon be realtime.





Rendering with distance fields

- Related info:

- “Making graphics in 4 kilobytes”: <http://www.rgb.org/iq/divulgation/inspire2008/inspire2008.htm>
- “Advanced perlin noise”: <http://rgba.scenes.org/iq/computer/articles/morenoise/morenoise.htm>